

2025-04-10 13:46

tags :

- [Machine Learning](#)
- [Hands on ML - book](#)

Chapter 3 - Hands on

• MNIST

- The `sklearn.datasets` package offers three main types of functions:
 - **fetch_*** **functions** (e.g., `fetch_openml()`): Download real-world datasets.
 - **load_*** **functions**: Load built-in small toy datasets (no internet needed).
 - **make_*** **functions**: Generate synthetic datasets for testing.
- Generated datasets usually come as `(X, y)` NumPy arrays, while others are returned as `Bunch` objects (dictionary-like), typically containing:
 - "DESCR" : Dataset description
 - "data" : Input features (2D array)
 - "target" : Labels (1D array)
- The `fetch_openml()` function typically returns inputs as a Pandas DataFrame and labels as a Series. However, for image datasets like MNIST, it's better to set `as_frame=False` to get NumPy arrays, which are more suitable for image data.

```
from sklearn.datasets import fetch_openml
mnist = fetch_openml('mnist_784', as_frame=False)
X, y = mnist.data, mnist.target # X.shape->(70000, 784) y.shape->
(70000, )
```

- Each of the 70,000 MNIST images has 784 features (28×28 pixels), where each feature represents a pixel's intensity from 0 (white) to 255 (black). To visualize an image, reshape its feature vector to 28×28 and display it using `matplotlib.pyplot.imshow()` with `cmap="binary"` for grayscale. Here's a sample function:

```
def plot_digit(image_data):
    image = image_data.reshape(28, 28)
    plt.imshow(image, cmap="binary")
    plt.axis("off")

some_digit = X[0]
```

```
plot_digit(some_digit)
plt.show()
```

Training a Binary Classifier

- To build a simple digit classifier, we focus on detecting the number **5**, turning it into a **binary classification** problem (5 vs. not 5). We create target labels using:
- Then, we use Scikit-Learn's `SGDClassifier`, which is **efficient for large datasets and supports online learning**:

```
from sklearn.linear_model import SGDClassifier
y_train_5 = (y_train == '5')
y_test_5 = (y_test == '5')

sgd_clf = SGDClassifier(random_state=42)
sgd_clf.fit(X_train, y_train_5)

sgd_clf.predict([some_digit]) # Output: array([True])
```

Performance Measures

Measuring Accuracy Using Cross-Validation

- To evaluate the **SGDClassifier**, we use **3-fold cross-validation** with `cross_val_score`, which gives over **95% accuracy**—seemingly impressive
- However, a **DummyClassifier** that always predicts "not 5" still achieves **over 90% accuracy**, because only ~10% of the data are actually 5s:

```
cross_val_score(sgd_clf, X_train, y_train_5, cv=3, scoring="accuracy")
# Output: array([0.95035, 0.96035, 0.9604])
```

- **Conclusion:** Accuracy can be misleading on **imbalanced datasets**. A better alternative for evaluation is the **confusion matrix**, which gives more insight into model performance.
- **IMPLEMENTING CROSS-VALIDATION**

- To manually implement **cross-validation** with more control, you can use `StratifiedKFold`, which ensures each fold has a representative class distribution Like CROSS-VALIDATION. Here's the process:
 1. Use `StratifiedKFold(n_splits=3)` to split the data into 3 balanced folds.
 2. For each fold:
 - Clone the classifier.

- Train on the training folds.
- Predict on the test fold.
- Calculate and print the accuracy.

This replicates what `cross_val_score()` does, but allows customization. Example output:

`0.95035 , 0.96035 , 0.9604` — matching the automated method.

```
from sklearn.model_selection import StratifiedKFold
from sklearn.base import clone
skfolds = StratifiedKFold(n_splits=3) # add shuffle=True if the dataset
is not already shuffled
for train_index, test_index in skfolds.split(X_train, y_train_5):
    clone_clf = clone(sgd_clf)
    X_train_folds = X_train[train_index]
    y_train_folds = y_train_5[train_index]
    X_test_fold = X_train[test_index]
    y_test_fold = y_train_5[test_index]
    clone_clf.fit(X_train_folds, y_train_folds)
    y_pred = clone_clf.predict(X_test_fold)
    n_correct = sum(y_pred == y_test_fold)
    print(n_correct / len(y_pred)) # prints 0.95035, 0.96035, and
0.9604
```

Confusion Matrices

- A **confusion matrix** shows how often instances of class **A** are predicted as class **B**.
- You use it to evaluate **classification performance** in more detail than accuracy.
- **Code to build it:**

```
from sklearn.model_selection import cross_val_predict
y_train_pred = cross_val_predict(sgd_clf, X_train, y_train_5, cv=3) #it is
1d array contain 1 value for each instance, we pass the label to train the
model

from sklearn.metrics import confusion_matrix
cm = confusion_matrix(y_train_5, y_train_pred)

[[53892  687] # Non-5s: 53892 correct (TN), 687 mistaken as 5s (FP)
 [1891  3530]] # 5s: 1891 missed (FN), 3530 correct (TP)
```

- **Interpretation:**

- Rows = actual class
- Columns = predicted class
- Diagonal = correct predictions

- Off-diagonal = errors (false positives/negatives)
- A **perfect classifier** would have only true positives and true negatives
- While the **confusion matrix** is detailed, you might want simpler metrics:
 - **Precision** (a.k.a. positive prediction accuracy):
 - Measures how **many predicted positives are actually correct**.
 - A model can get **100% precision** by making **very few positive predictions**—but it may miss many actual positives.
 - **Recall** (a.k.a. sensitivity or true positive rate):
 - Measures how **many actual positives the model correctly detects**.
 - High recall means **fewer false negatives**.

Precision and Recall

- After evaluating the 5-detector using accuracy, it's time to dig deeper using **precision**, **recall**, and the **F1 score**.
- **Metrics:**
 - **Precision:**
 - → Of all predicted 5s, **83.7% were correct**.
 - **Recall:**
 - → The model found **65.1% of actual 5s**.
- **F1 Score** (harmonic mean of precision and recall):
 - Balances both precision and recall.
 - Favors models with **similar** values for both.
- You **can't maximize both** at the same time.
 - Choose based on context:
 - **High precision, low recall:** Good for **kid-safe content filters**.
 - **High recall, low precision:** Okay for **shoplifter detection**, where missing one is worse than a false alarm.

The Precision/Recall Trade-off

- **SGDClassifier** computes a **decision score** for each input.
- If the score is **above a threshold**, it predicts **positive** (e.g., “5”); otherwise, **negative**.
- **Raising the threshold:**
 - Increases **precision** (fewer false positives)

- ✗ Decreases **recall** (more false negatives)
- **Lowering the threshold:**
 - ✓ Increases **recall**
 - ✗ Decreases **precision**
- **Example:**
 - If threshold is at default 0 : prediction is **True** (high recall).
 - Raising it to 3000 : prediction becomes **False** (misses a correct 5 → lower recall).
- **How to Tune the Threshold:**

1. **Get decision scores** (instead of predictions):

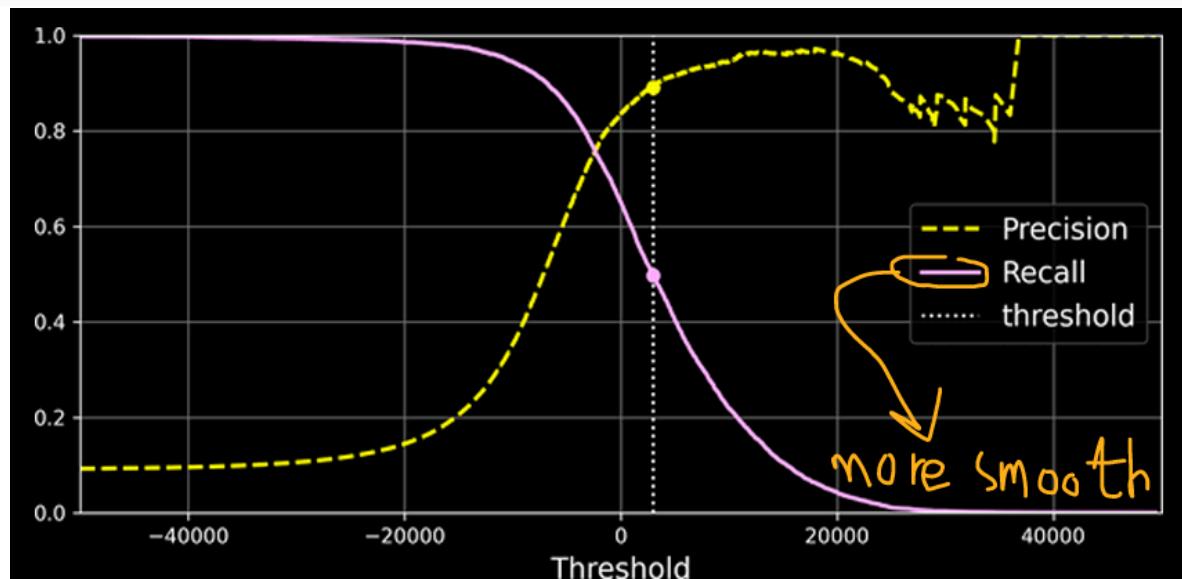
```
y_scores = cross_val_predict(sgd_clf, X_train, y_train_5, cv=3,
                             method="decision_function")
```

2. **Compute precision and recall for all thresholds:**

```
from sklearn.metrics import precision_recall_curve
precisions, recalls, thresholds = precision_recall_curve(y_train_5,
                                                          y_scores)
```

3. **Visualize precision & recall against thresholds:**

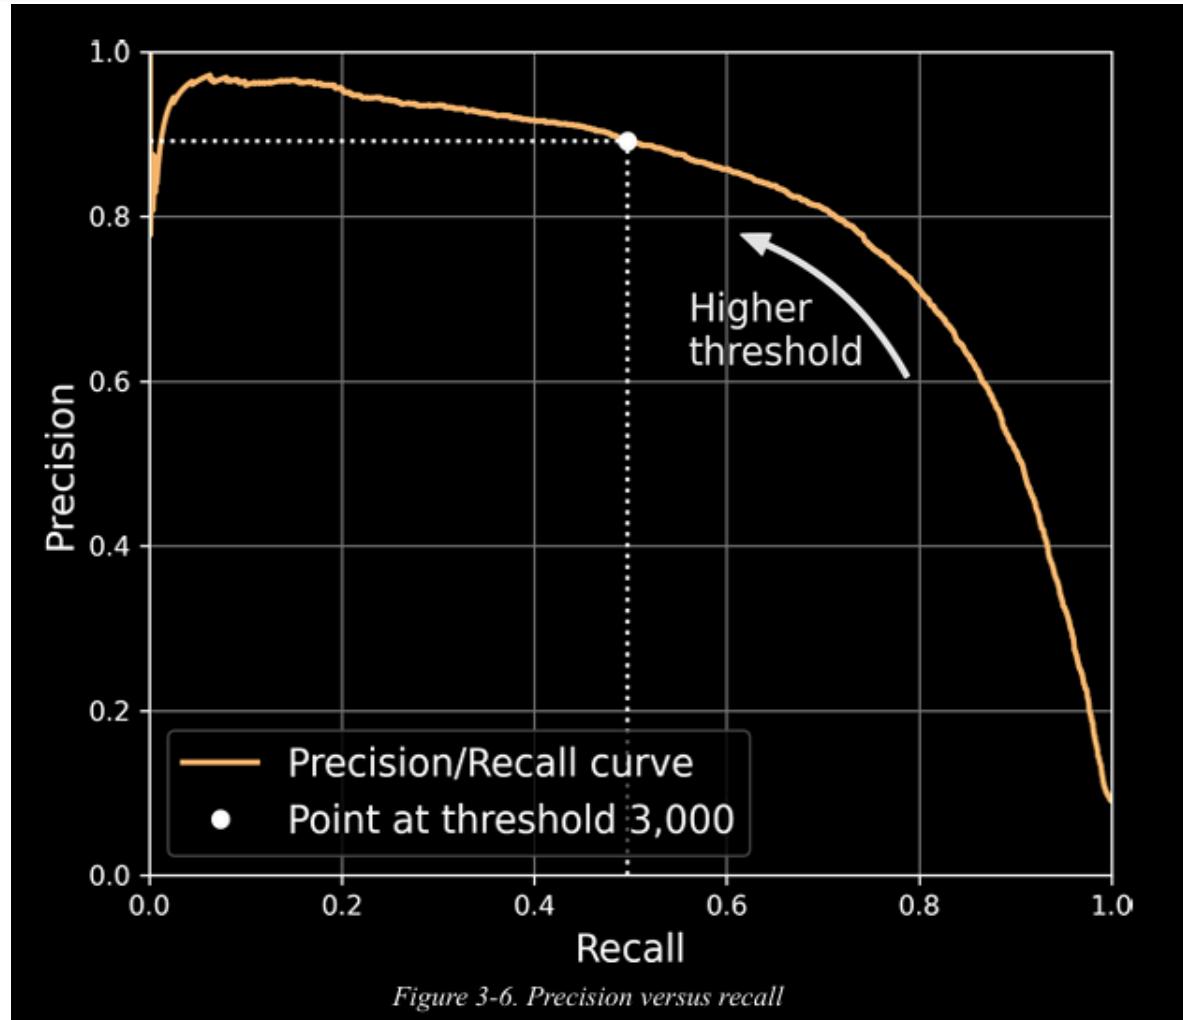
```
plt.plot(thresholds, precisions[:-1], "b--", label="Precision")
plt.plot(thresholds, recalls[:-1], "g-", label="Recall")
# note here we exclude the last value in the precisions and recalls
# because recalls nad precisions have length N+1 but thresholds have
# length = N
plt.vlines(threshold, 0, 1.0, "k", "dotted", label="threshold")
# plt.vlines(x, ymin, ymax, color, linestyle, label)
```



• *Figure 3-5. Precision and recall versus the decision threshold*

- 4. Or plot precision directly vs. recall to spot the sweet spot before a sharp drop in precision.

```
plt.plot(recalls, precisions, linewidth=2, label="Precision/Recall curve")
#[...] beautify the figure: add labels, grid, legend, arrow, and text
plt.show()
```



- Key Insight:
 - There's **no universal best threshold**.
 - Choose it based on **what matters more**:
 - **Precision**: when **false positives are costly** (e.g., kid-safe content).
 - **Recall**: when **missing positives is worse** (e.g., medical diagnosis or shoplifter detection).
- You can **tune your classifier** to reach **90% precision** by:

```
idx = (precisions >= 0.90).argmax()
threshold = thresholds[idx]
y_pred_90 = (y_scores >= threshold)
# hecking performance Precision: 90% and Recall: Only ~48%
```

The ROC Curve

- It plots **True Positive Rate (TPR) = Recall** vs **False Positive Rate (FPR)**.
- $FPR = 1 - \text{Specificity} = \text{fraction of negatives misclassified as positives}$.
- How is it constructed? (code)

```
from sklearn.metrics import roc_curve
fpr, tpr, thresholds = roc_curve(y_train_5, y_scores)

#plotting
idx_for_threshold_at_90 = (thresholds <=
threshold_for_90_precision).argmax()
tpr_90, fpr_90 = tpr[idx_for_threshold_at_90], fpr[idx_for_threshold_at_90]
plt.plot(fpr, tpr, linewidth=2, label="ROC curve") plt.plot([0, 1], [0, 1],
'k:', label="Random classifier's ROC curve")
plt.plot([fpr_90], [tpr_90], "ko", label="Threshold for 90% precision")
# [...] beautify the figure: add labels, grid, legend, arrow, and text
plt.show()
```

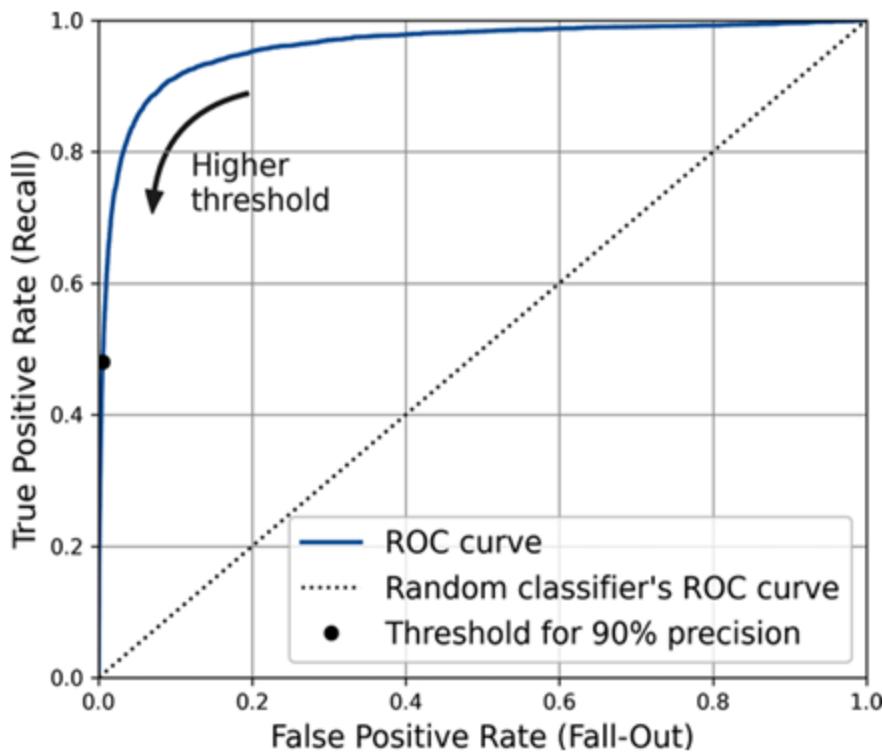


Figure 3-7. A ROC curve plotting the false positive rate against the true positive rate for all possible thresholds; the black circle highlights the chosen ratio (at 90% precision and 48% recall)

- **A great classifier** hugs the **top-left corner** (high recall, low FPR).
- **The farther above the diagonal** it is, the better.
- AUC (Area Under the Curve):

```
from sklearn.metrics import roc_auc_score
```

```
roc_auc_score(y_train_5, y_scores)
```

- Value close to **1.0** = excellent classifier.
- Value **0.5** = random guessing.

- **Choosing Between ROC and PR Curves**

- Use **PR curve** when:
 - The **positive class is rare**.
 - You care **more about false positives** than false negatives.
- Use **ROC curve** otherwise.
- ROC AUC can be **misleading** if positives are rare—it might look good even if the model isn't.
- PR curve gives a more honest view in such cases.

- **Random Forest vs SGD Classifier (on PR Curve)**

1. **RandomForestClassifier**:

- **Doesn't** use `decision_function()` like SGD.
- Use `predict_proba()` to get probability scores instead.
- Use the **positive class probability** as the score for PR/ROC analysis.

```
from sklearn.ensemble import RandomForestClassifier
forest_clf = RandomForestClassifier(random_state=42)

y_probas_forest = cross_val_predict(forest_clf, X_train, y_train_5, cv=3,
method="predict_proba")

y_probas_forest[:2]
# array([[0.11, 0.89],
#        [0.99, 0.01]]) First image: 89% likely to be positive, Second
image: 99% likely to be negative.
```

- **Warning About Probabilities**

- These are **estimated probabilities**, not guaranteed.
- For example, samples predicted with 50–60% confidence were actually positive **94% of the time** → model underestimated.
- Some models can also be **overconfident**.
- Use `sklearn.calibration` to fix this and make probabilities more accurate.

- **Evaluating with PR Curve**

- Use `precision_recall_curve()` on `y_scores_forest = y_probas_forest[:, 1]`.
- Plot both **Random Forest** and **SGD** PR curves for comparison.

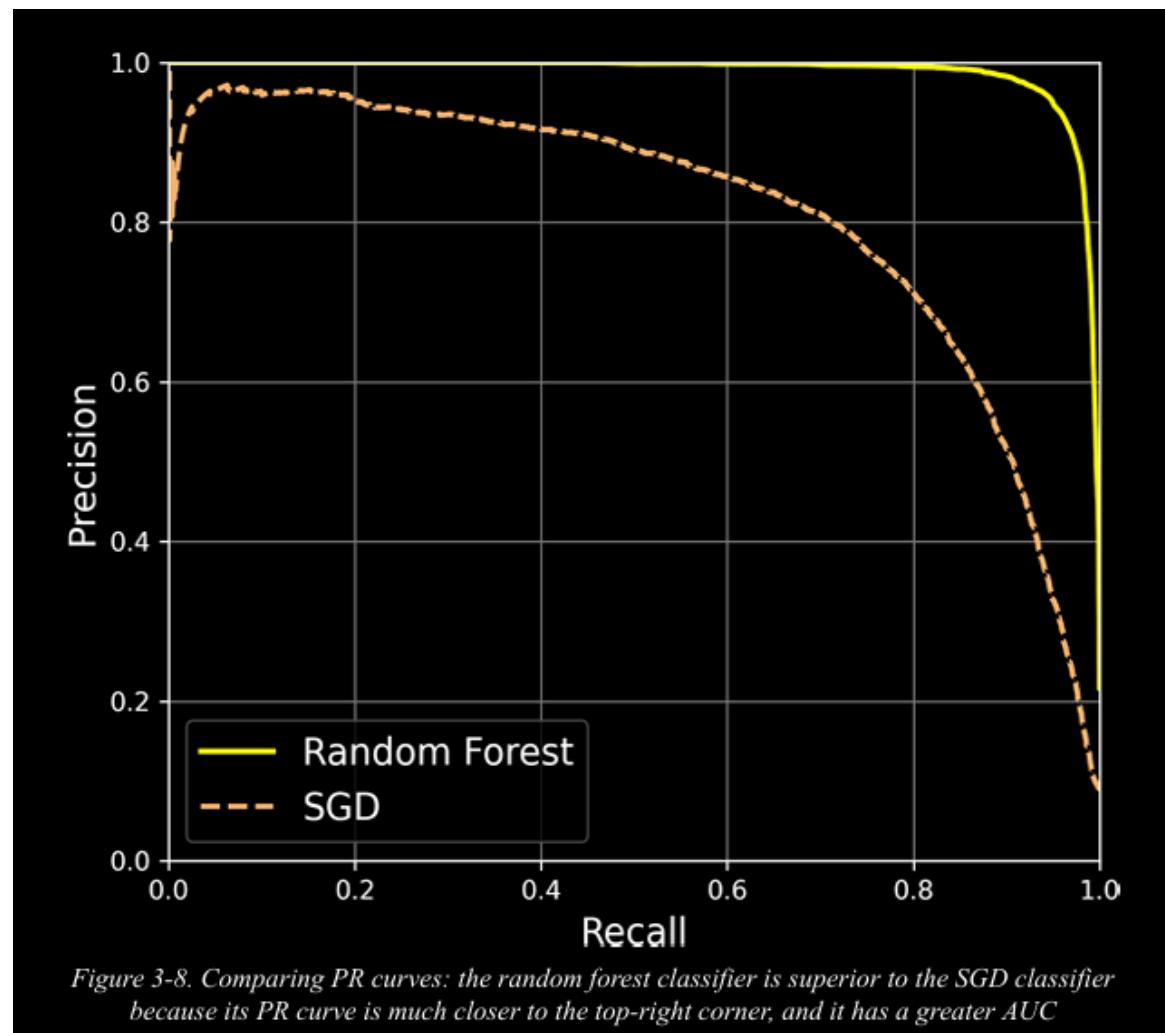


Figure 3-8. Comparing PR curves: the random forest classifier is superior to the SGD classifier because its PR curve is much closer to the top-right corner, and it has a greater AUC

- Random Forest performs much better:
 - Closer PR curve to the top-right corner.
 - **F1 score:** 0.924
 - **ROC AUC score:** 0.998
 - **Precision** ≈ 99.1%, **Recall** ≈ 86.6%

Multiclass Classification

- What is Multiclass Classification?
 - Unlike binary classification (2 classes), **multiclass classification** handles **more than two** classes (e.g., digits 0–9).
- Built-in Support
 - Some classifiers (like `LogisticRegression`, `RandomForestClassifier`) support multiclass natively.
 - Others (like `SGDClassifier`, `SVC`) only support binary classification—but **Scikit-Learn** handles multiclass automatically using two strategies:
- Strategies for Multiclass with Binary Classifiers
 1. **One-vs-Rest (OvR / OvA)**

- Train **one binary classifier per class**.
- Classify by picking the classifier with the **highest decision score**.

2. One-vs-One (OvO)

- Train a binary classifier for **every pair of classes**.
- For 10 classes (MNIST): **45 classifiers** → $10 \times 9 / 2$.
- Each classifier “votes”; the class with **most wins** is chosen.
- The main advantage of OvO is that each classifier only needs to be trained on the part of the training set containing the two classes that it must distinguish.

- Example: SVC with OvO

- SVC scales poorly with large datasets, so **OvO is used by default**.

```
svm_clf = SVC(random_state=42)
svm_clf.fit(X_train[:2000], y_train[:2000])
```

- It trains **45 classifiers** internally for 10 classes.

- When you predict, it:

- Runs the sample through all 45 classifiers.
- Collects **10 class scores** using `decision_function()` — one score per class (based on votes).
- Picks the class with the **highest score**.

```
some_digit_scores = svm_clf.decision_function([some_digit])
class_id = some_digit_scores.argmax()
svm_clf.classes_[class_id] # gives predicted class
```

- Even though OvO uses many classifiers, **you still get one score per class**.

- You can **manually choose** One-vs-One (OvO) or One-vs-Rest (OvR) using:

- `OneVsOneClassifier`
- `OneVsRestClassifier`

```
from sklearn.multiclass import OneVsRestClassifier
ovr_clf = OneVsRestClassifier(SVC(random_state=42))
ovr_clf.fit(X_train[:2000], y_train[:2000]) #This trains **10 classifiers**
#one for each class.
ovr_clf.predict([some_digit]) # Output: '5'
```

- Using SGDClassifier for Multiclass

```
sgd_clf = SGDClassifier(random_state=42)
sgd_clf.fit(X_train, y_train)
sgd_clf.predict([some_digit]) # Output: '3' (could be wrong)
```

```
cross_val_score(sgd_clf, X_train, y_train, cv=3, scoring="accuracy")
# → ~86% accuracy (better than random guessing, which is 10%)
```

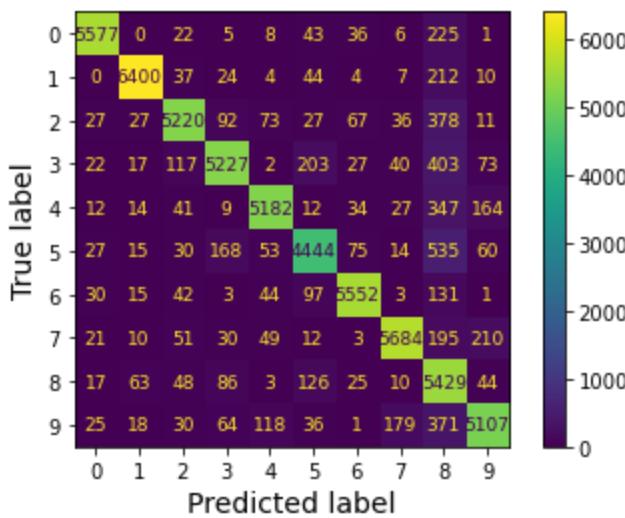
- SGDClassifier uses **OvR** by default.
- decision_function([some_digit]) returns **10 confidence scores**, one per class.
- Boosting Accuracy with Scaling:

```
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train.astype("float64"))
cross_val_score(sgd_clf, X_train_scaled, y_train, cv=3,
scoring="accuracy")
# → ~89–90% accuracy
```

Error Analysis

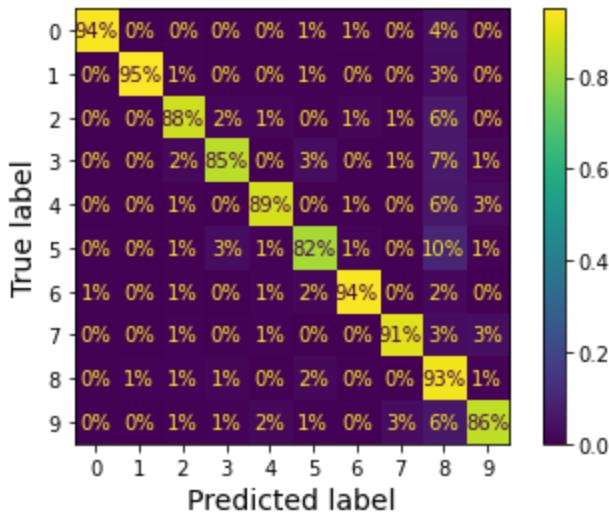
- Once you've trained a promising model, analyze **what types of mistakes it makes** to improve it.
- Confusion Matrix
 - Use cross_val_predict() to get predictions.
 - Plot the confusion matrix using:

```
from sklearn.metrics import ConfusionMatrixDisplay
y_train_pred = cross_val_predict(sgd_clf, X_train_scaled, y_train, cv=3)
ConfusionMatrixDisplay.from_predictions(y_train, y_train_pred)
plt.show()
```



- Normalized Confusion Matrix

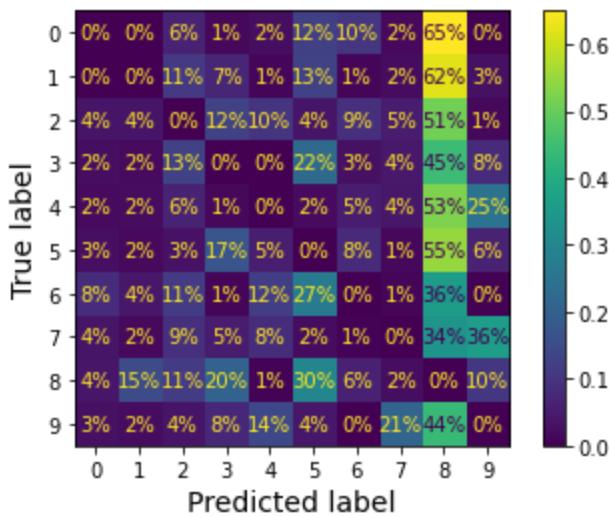
```
ConfusionMatrixDisplay.from_predictions(y_train, y_train_pred,
normalize="true",
values_format=".0%")
```



- Example finding: Only 82% of digit 5s are correctly predicted.
 - 10% of 5s were misclassified as 8s.
 - Only 2% of 8s were misclassified as 5s → **asymmetrical errors**.

- Highlight Only the Errors(To make misclassifications more visible)

```
sample_weight = (y_train_pred != y_train)
ConfusionMatrixDisplay.from_predictions(y_train, y_train_pred,
                                        sample_weight=sample_weight,
                                        normalize="true",
                                        values_format=".0%")
```



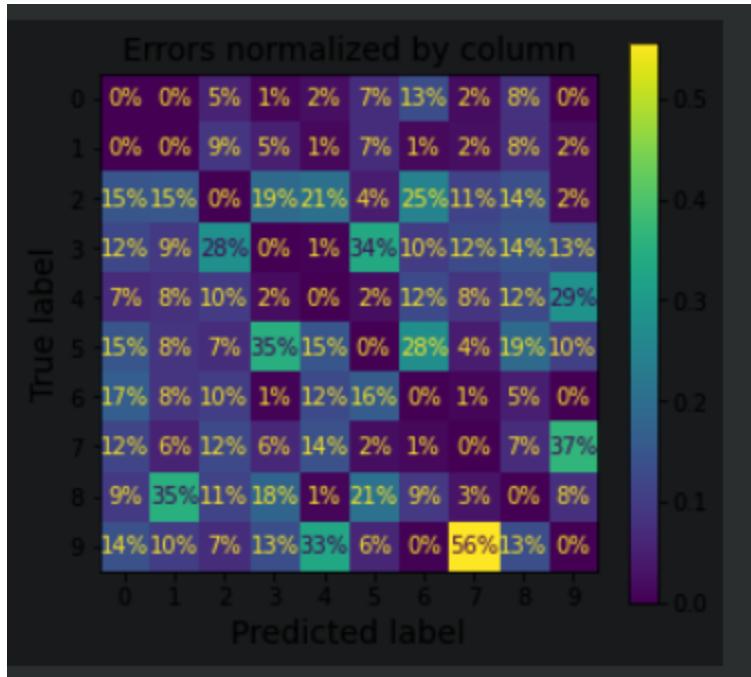
- Column-Normalized Confusion Matrix
 - Use `normalize="pred"` to normalize by column:
 - Shows **what predicted class the errors are going to**.
 - Example: 56% of predicted 7s are actually 9s.

```
ConfusionMatrixDisplay.from_predictions(
    y_train, y_train_pred,
    ax=axs[1],
```

```

        sample_weight=sample_weight,
        normalize="pred", # the change
        values_format=".0%"
    )

```



- Error Analysis:
 - Use a **confusion matrix** to find common mistakes (e.g., many digits misclassified as 8).
 - Improve accuracy by:
 - Adding **more training data** for confusing digits.
 - **Engineering features** (e.g., count loops in digits).
 - **Preprocessing images** to highlight important patterns.
 - **Visualize errors** (like 3s vs 5s) to see where and why mistakes happen.
 - **SGDClassifier is linear**, so it's sensitive to small changes like shifts or rotations.
 - **Best fix:** Use **data augmentation**—add rotated/shifted versions of training images to improve model robustness.

Multilabel Classification

- In **multilabel classification**, each input can be assigned **multiple binary labels** (not just one class).

```

y_train_large = (y_train >= '7')           # True if digit is 7, 8, or 9
y_train_odd = (y_train.astype('int8') % 2 == 1) # True if digit is odd
y_multilabel = np.c_[y_train_large, y_train_odd]

```

- This creates **two labels per digit**:
- Then trains a **KNeighborsClassifier** with these multilabels.

- Prediction Example

```
knn_clf.predict([some_digit]) → [[False, True]]
```

- → The digit is not large but it is odd (like 5).

- Evaluation

- Use **F1 score** with `average="macro"` to give equal weight to each label:

```
f1_score(y_multilabel, y_train_knn_pred, average="macro")
```

- → `average="weighted"` gives more weight to labels with more examples.

- When Classifier Doesn't Support Multilabels

- Train one model per label (but they won't communicate).
- A better option: use `ClassifierChain`.

- `ClassifierChain`

- Organizes models **in a chain** so later models use predictions from earlier ones.
- Use `cv` to get clean predictions during training.

```
chain_clf = ClassifierChain(SVC(), cv=3, random_state=42)
chain_clf.fit(X_train[:2000], y_multilabel[:2000])
chain_clf.predict([some_digit]) # → [[0., 1.]]
```

Multioutput Classification

- What is Multioutput Classification?

- A generalization of multilabel classification.
- Each label can have **more than two values** (i.e., it's multiclass per output).
- Example: A system that removes noise from images.

- Noise Removal Example

- Input: **Noisy MNIST digit** image (added random noise to each pixel).
- Output: **Clean digit** image (original pixel intensities).
- Each pixel is a label, and its value ranges from 0–255 → **multioutput multiclass**.

- code

```
noise = np.random.randint(0, 100, (len(X_train), 784))
X_train_mod = X_train + noise          # Noisy inputs
y_train_mod = X_train                  # Clean targets
```

- Trained a **KNeighborsClassifier** to predict clean images from noisy ones:

```
knn_clf.fit(X_train_mod, y_train_mod)
clean_digit = knn_clf.predict([X_test_mod[0]])
```

Resources :

•

Related notes :

•

References :

- **Internal :**

-

-

-

- **External :**

- [notebook of the chapter](#)

- [hegab videos](#)

- [the book](#)

-