

2025-04-28 00:43

## tags :

- [Machine Learning](#)
- [Hands on ML - book](#)

## Chapter 5 - Hands on

- **SVMs** are powerful models for **classification**, **regression**, and **novelty detection**.
- They work best on **small to medium-sized**, **nonlinear** datasets but don't scale well to **very large datasets**.

## Linear SVM Classification

- **Goal**: Find a straight line (decision boundary) that separates two classes **with the largest margin**.
- **Large / hard Margin Classification**: SVM chooses the line that leaves the **widest possible street** between classes.
- **Support Vectors**: Only the instances on the edge of the street (circled points) determine the decision boundary.
- important Warning
  - **SVMs are sensitive to feature scales**.
  - If features aren't scaled (e.g., vertical much bigger than horizontal), the SVM decision boundary can be wrong.
  - After **feature scaling** (like using `StandardScaler`), the boundary becomes better and more balanced. illustration [\[1\]](#)

## Soft Margin Classification

- **Hard Margin Classification** forces all points to stay outside the margin — works only if data is perfectly separable and is **very sensitive to outliers**.
- **Soft Margin Classification** allows some violations (some points inside the margin or misclassified) to balance between **margin size** and **errors** for better generalization.
- **C Hyperparameter** controls the trade-off:
  - **Small C** → Larger margin, more violations → Less overfitting but risk of underfitting.
  - **Large C** → Smaller margin, fewer violations → Risk of overfitting but fits training data better.

- **Tip:** If the SVM overfits, **reduce C** to regularize.
- **Code Example:** Using `LinearSVC` inside a `Pipeline` (with `StandardScaler`) to classify Iris Virginica flowers.

```
from sklearn.datasets import load_iris
from sklearn.pipeline import make_pipeline
from sklearn.preprocessing import StandardScaler
from sklearn.svm import LinearSVC

iris = load_iris(as_frame=True)
X = iris.data[["petal length (cm)", "petal width (cm)"]].values
y = (iris.target == 2) # Iris virginica
svm_clf = make_pipeline(StandardScaler(),
                        LinearSVC(C=1, random_state=42))

svm_clf.fit(X, y)

X_new = [[5.5, 1.7], [5.0, 1.5]]
svm_clf.predict(X_new) # O/P => array([ True, False])
```

- **Decision Function:** Outputs a score showing how far each instance is from the decision boundary.


```
svm_clf.decision_function(X_new) # O/P => array([ 0.66163411,
        -0.22036063])
```

- **Note:** `LinearSVC` doesn't support probability predictions, but `SVC(probability=True)`
  - Normally, `LinearSVC` only gives a **score**, not a **probability**.
  - If you use `SVC(probability=True)`:
    - The model learns to **map scores to probabilities** (e.g., 80% confidence).
    - To do that, it must:
      - Run **5-fold cross-validation** during training.
      - Train a **Logistic Regression model** on the cross-validation results to map scores to probabilities.
  - ➡ **This extra work makes training slower.**

## Nonlinear SVM Classification

- Many datasets are **not linearly separable**.
- You can **add new features** (like polynomials) to make the data **linearly separable** in a higher-dimensional space.

- Example: Adding a second feature like  $x_2 = x_1^2$  makes a previously non-separable dataset separable.



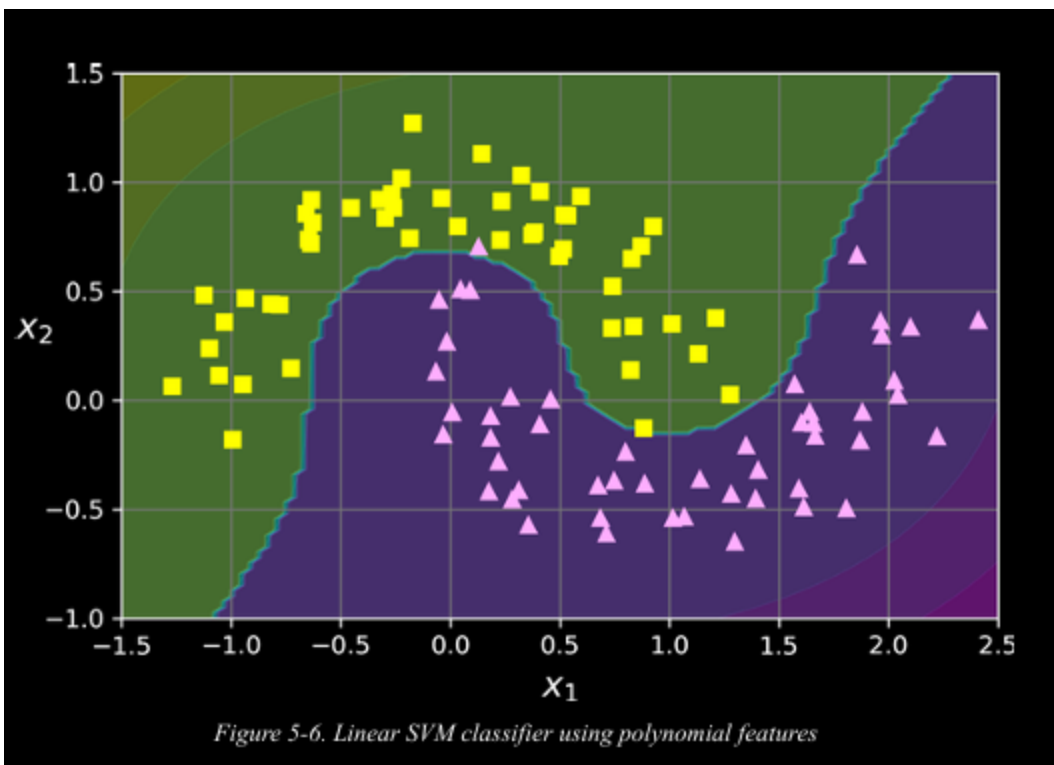
- **How to do it:**
  - Use `PolynomialFeatures` to create new features.
  - Then scale features with `StandardScaler`.
  - Then train a `LinearSVC` model.
  - Example code:

```
from sklearn.datasets import make_moons
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import PolynomialFeatures, StandardScaler
from sklearn.svm import LinearSVC

X, y = make_moons(n_samples=100, noise=0.15, random_state=42)

polynomial_svm_clf = Pipeline([
    ("poly_features", PolynomialFeatures(degree=3)),
    ("scaler", StandardScaler()),
    ("svm_clf", LinearSVC(C=10, loss="hinge", random_state=42))
])

polynomial_svm_clf.fit(X, y)
```



## Polynomial Kernel

- Adding polynomial features can help models handle nonlinearity, but:
  - Low degrees can't model complex patterns well.
  - High degrees cause **too many features** → slow models.
- **Kernel trick**: SVMs can act like they added high-degree polynomial features **without** actually creating them — no feature explosion!
- Use `SVC` with `kernel="poly"` to apply the polynomial kernel easily.
- Example:

```
from sklearn.svm import SVC
from sklearn.pipeline import make_pipeline
from sklearn.preprocessing import StandardScaler


poly_kernel_svm_clf = make_pipeline(
    StandardScaler(),
    SVC(kernel="poly", degree=3, coef0=1, C=5)
)
poly_kernel_svm_clf.fit(X, y)
```

- **Important hyperparameters**:
  - `degree` : degree of the polynomial (higher = more flexible, but risk of overfitting).
  - `coef0` : controls balance between high-degree and low-degree features.


- **Tips:**
  - If overfitting → **lower** the degree.
  - If underfitting → **raise** the degree.
- Even when hyperparameters are tuned automatically, it's important to **understand them** to limit the search space. use randomized search

## Similarity Features

- Another way to handle **nonlinear data** is by **adding features based on similarity** to certain **landmarks**.
- Example:
  - Landmarks placed at  $x=-2$  and  $x=1$ .
  - Use **Gaussian RBF** (Radial Basis Function) with  $\gamma=0.3$  to measure similarity.
  - Similarity is high (near 1) if close to a landmark, and low (near 0) if far.



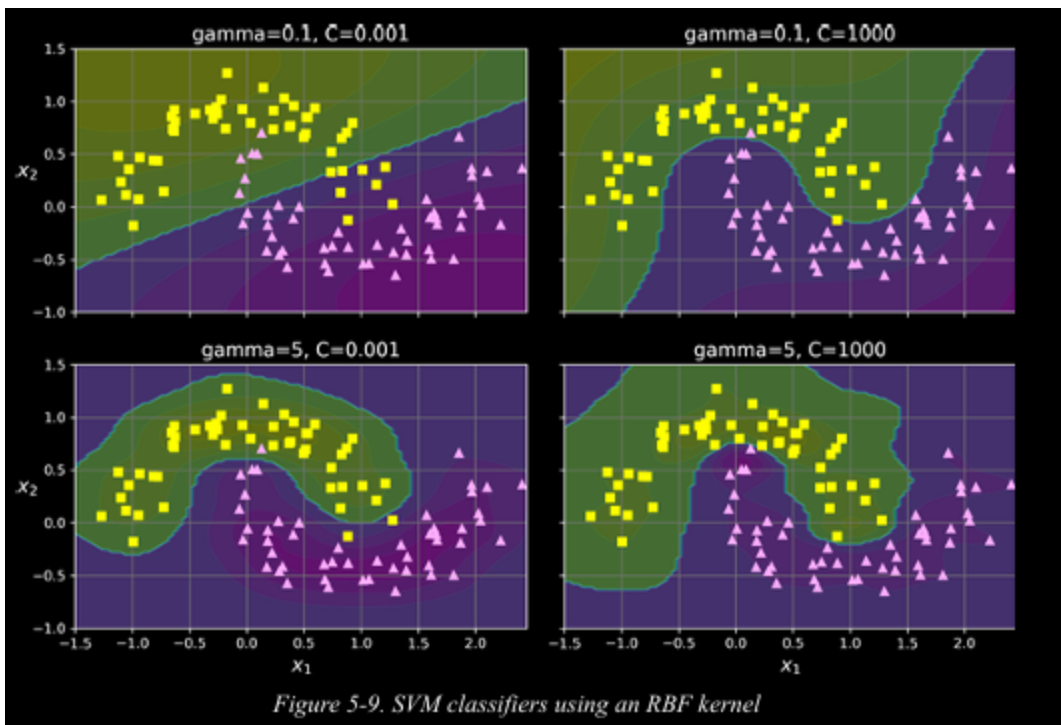
- For instance  $x=-1$ :
  - Similarity to first landmark: about **0.74**.
  - Similarity to second landmark: about **0.30**.
- After transformation, the dataset can become **linearly separable**.



- **Landmark selection:**
  - Easiest way: make every training instance a landmark.
  - Downside: if you have **m instances**, you get **m features** — can be very **computationally heavy** for large datasets.

# Gaussian RBF Kernel

- Like polynomial features, similarity features can be useful but **computationally heavy**.
- The **kernel trick** in SVMs allows getting the same effect **without** explicitly adding features.
- Example: using `SVC(kernel="rbf", gamma=5, C=0.001)`.
- **Gamma ( $\gamma$ )**:
  - Higher  $\gamma$  → narrower bell curve → smaller influence → **more irregular (overfitting)** decision boundary.
  - Lower  $\gamma$  → wider bell curve → larger influence → **smoother (underfitting)** decision boundary.



- $\gamma$  acts like a **regularization hyperparameter** (similar to  $C$ ).
- Other kernels exist (e.g., for text or DNA sequences) but are less common.
- **Tip**:
  - Try a **linear kernel first** (preferably `LinearSVC` for speed).
  - If the dataset is small enough, try **Gaussian RBF kernel** next.
  - Experiment with other specialized kernels if needed.

## SVM Classes and Computational Complexity

- **LinearSVC**:
  - Based on `liblinear`.
  - **No kernel trick**.
  - Fast:  $O(m \times n)$  time complexity.
  - Needs scaling.

- No out-of-core learning. explain -> [2]
- **SVC:**
  - Based on **libsvm**.
  - **Supports kernel trick** (for nonlinear problems).
  - **Slow** on large datasets:  $O(m^2 \times n)$  to  $O(m^3 \times n)$ .
  - Needs scaling.
  - No out-of-core learning.
- **SGDClassifier:**
  - Uses **stochastic gradient descent**.
  - $O(m \times n)$  time complexity.
  - **Supports out-of-core learning** (good for huge datasets).
  - Needs scaling.
  - No kernel trick.

## SVM Regression

- **SVM for Regression:**
  - Instead of separating classes, SVM tries to fit as many points **inside a margin ("street")** as possible.
  - The **width of the street** is controlled by **epsilon ( $\epsilon$ )**.
  - Points inside the margin **don't affect** the model → called  **$\epsilon$ -insensitive**.
- **LinearSVR:** `from sklearn.svm import LinearSVR`
  - For **linear regression** (like a straight line).
  - Fast and scales well to large datasets.
- **SVR** (with kernel): `from sklearn.svm import SVR`
  - For **nonlinear regression** (curves, etc.).
  - Slow on large datasets because it uses the **kernel trick**.
- **C hyperparameter:**
  - Controls how **strict** the model is about margin violations.
  - **Small C** → more flexible (allow violations, more regularization).
  - **Large C** → less flexible (less violations, tries to fit the training data closely).
- **epsilon ( $\epsilon$ ) hyperparameter:**
  - Controls the **width of the margin ("street")** where no penalty is given for errors.
  - **Small  $\epsilon$**  → narrow margin → **more support vectors** → **more complex model**.
  - **Large  $\epsilon$**  → wide margin → simpler model.

Concept	C (penalty)	epsilon (street width)
Controls	How much you allow errors outside margin	How wide the no-penalty zone is
Small value	More regularization, allow more errors	Narrower street, more support vectors
Large value	Fit more tightly to training data	Wider street, fewer support vectors

## Under the Hood of Linear SVM Classifiers

- A **linear SVM classifier** predicts the class of an input by computing a **decision function**:  
 $w^\top x + b$ 
  - If the result is **positive**, it predicts class **1**.
  - If the result is **negative**, it predicts class **0**.
- This is similar to **Logistic Regression** in prediction style.
- **Two notations for parameters**:
  - **Old way**: Combine weights and bias into one vector  $\theta$  (with a fake feature  $x_0 = 1$ ).
  - **New way (standard)**:
    - **w** = weights vector
    - **b** = bias (separate)
 So, prediction =  $w^\top x + b$ .
- **Training a Linear SVM**:
  - **Goal**:
    - **Maximize the margin** (make the street between classes as wide as possible).
    - **Minimize margin violations** (keep instances outside the street as few as possible).
  - **Key ideas**:
    - **Smaller w**  $\rightarrow$  **larger margin** (margin  $\propto 1/\|w\|$ ).
    - **Bias b** only **shifts** the street but **does not change** its width.
- **Optimization Problem (Hard Margin SVM)**:
  - Minimize:
 
$$\frac{1}{2} w^\top w$$
  - Subject to:
 
$$t^{(i)}(w^\top x^{(i)} + b) \geq 1 \quad \text{for all instances}$$
 where  $t^{(i)} = 1$  if positive class,  $-1$  if negative class.
- We minimize  $\frac{1}{2} w^\top w$  instead of  $\|w\|$  because:
  - It's **differentiable** (makes optimization easier).



- Derivative of  $\frac{1}{2}w^\top w$  is simply  $w$ .


Concept	Meaning
$w$ small	Larger margin (better separation)
$b$	Shifts the decision boundary (no size change)
Objective	Minimize $\frac{1}{2}\ w\ ^2$ while keeping predictions correct
Prediction Rule	Positive if $w^\top x + b > 0$ , negative otherwise

- In real-world data, **perfect separation** may be impossible.
- To handle this, we introduce **slack variables**  $\zeta^{(i)} \geq 0$  for each instance:
  - Each  $\zeta$  measures **how much margin violation** is allowed for that instance.
- Soft Margin SVM Objective:**
  - We want to **Minimize** both:
    - The model complexity ( $\frac{1}{2}\|w\|^2 \rightarrow$  wider margin)
    - The total margin violations (sum of  $\zeta$ 's)
  - C hyperparameter** controls the **trade-off** between margin size and violations:
    - High **C**  $\rightarrow$  less tolerance for margin violations (harder margin).
    - Low **C**  $\rightarrow$  more tolerance (softer margin).
  - The optimization problem becomes: Minimize  $\frac{1}{2}w^\top w + C \sum_{i=1}^m \zeta^{(i)}$
  - subject to:  $t^{(i)}(w^\top x^{(i)} + b) \geq 1 - \zeta^{(i)}, \quad \zeta^{(i)} \geq 0$
- Optimization Techniques:**
  - These are **convex quadratic programming (QP)** problems (special type of optimization with linear constraints).
  - Training options:**
    - Use a **QP solver** (specialized software).
    - Use **gradient descent** by minimizing:
      - Hinge loss** (linear penalty)
      - Squared hinge loss** (quadratic penalty, more sensitive to outliers).
- Hinge Loss Behavior:**

	Positive class (t=1)	Negative class (t=-1)
Correctly classified & far from margin	Loss = 0	Loss = 0
Inside margin or wrong side	Loss > 0	Loss > 0

- Hinge loss** grows **linearly** with error.

- **Squared hinge loss** grows **quadratically** with error (faster convergence if clean data).



- In Scikit-Learn:
  - **LinearSVC**: uses **squared hinge loss** by default.
  - **SGDClassifier**: uses **hinge loss** by default.
  - **SVC**: similar to minimizing hinge loss but with kernels.

Name	What it Is	Controls	Used in
<b>C</b> (capital letter)	Hyperparameter chosen by you	Trade-off between margin size and margin violations	Classification & Regression
<b>ε</b> (epsilon)	Width of the “street” (margin) where no penalty happens in <b>regression</b>	Tolerance for errors inside the margin (in regression)	<b>SVM Regression (SVR)</b>
<b>ζ</b> (zeta)	Slack variables computed during training	How much each individual training point <b>violates</b> the margin (in classification)	<b>SVM Classification (Soft margin)</b>

## The Dual Problem

→ What is happening?

- **SVM training** can be done using the **primal problem** or the **dual problem**.
- In SVM, the **dual** and **primal** give **the same solution** (because of special conditions).
- The **dual problem** is often easier when:
  - The number of training samples **m** is **less than** the number of features.
  - You want to use the **kernel trick** later.
- Dual Problem Equation  
We **\*\*solve for α** by minimizing:

$$\text{minimize} \quad \frac{1}{2} \sum_{i=1}^m \sum_{j=1}^m \alpha^{(i)} \alpha^{(j)} t^{(i)} t^{(j)} \left( x^{(i)} \cdot x^{(j)} \right) - \sum_{i=1}^m \alpha^{(i)}$$

subject to:  $\alpha^{(i)} \geq 0$  for all  $i = 1, 2, \dots, m$

$$\sum_{i=1}^m \alpha^{(i)} t^{(i)} = 0$$

- $\alpha^{(i)}$  = slack related values (Lagrange multipliers).
- $t^{(i)}$  = class label (+1 or -1).
- $x^{(i)}$  = feature vectors.
- From Dual Solution to Primal Solution

Once you find the optimal  $\hat{\alpha}$ , you can **rebuild** your model:

- **Weights vector**  $\hat{w}$  :

$$\hat{w} = \sum_{i=1}^m \hat{\alpha}^{(i)} t^{(i)} x^{(i)}$$

- **Bias term**  $\hat{b}$ :

$$\hat{b} = \frac{1}{n_s} \sum_{\text{support vectors}} \left( t^{(i)} - \hat{w}^\top x^{(i)} \right)$$

- where:
  - $n_s$  = number of **support vectors** (i.e., data points where  $\hat{\alpha}^{(i)} > 0$ ).
- Quick notes:
  - **Support vectors** are the important data points that "touch" the margin.
  - If the dataset is very big, **dual** can be slower (but necessary for kernels).
  - **Primal** is faster for very large datasets with simple (linear) relationships.

## Kernelized SVMs

- Problem:
  - You want to apply a **second-degree polynomial transformation**  $\phi(x)$  to your 2D data to make it easier to separate with a **linear SVM**.

Transformation (Equation 5-5):

$$\phi(\mathbf{x}) = \begin{pmatrix} x_1^2 \\ \sqrt{2}x_1x_2 \\ x_2^2 \end{pmatrix}$$

After transformation, the data becomes **3D** instead of 2D.

- Kernel Trick:
  - Instead of explicitly computing  $\phi(x)\phi(x)$ , notice:  
 $\phi(a)^\top \phi(b) = (a^\top b)^2$

**Key insight:**

- You can **just compute**  $(a^\top b)^2$  — **without ever computing**  $\phi(x)$  **directly!**  
(Thus saving a lot of computation.)
- Common Kernels (Equation 5-7):
  - **Linear kernel:**  $K(a, b) = a^\top b$
  - **Polynomial kernel:**  $K(a, b) = (\gamma a^\top b + r)^d$
  - **Gaussian RBF kernel:**  $K(a, b) = \exp(-\gamma \|a - b\|^2)$
  - **Sigmoid kernel:**  $K(a, b) = \tanh(\gamma a^\top b + r)$
- Mercer's Theorem:
  - If  $K(a, b)$  satisfies some properties (continuous, symmetric, positive semi-definite), then there exists a mapping  $\phi$  such that:  $K(a, b) = \phi(a)^\top \phi(b)$   
This guarantees that you can safely use kernels **without needing to know**  $\phi$ !
  - **Example:**  
For Gaussian RBF,  $\phi(x)$  is infinite-dimensional — but we never actually compute it.
  - **Note:**  
Some kernels (like sigmoid) **violate Mercer's conditions**, but **work well in practice**.
- Predictions with Kernelized SVMs:
  - When you train with kernels, you **can't compute**  $\hat{w}$  explicitly anymore (because  $\phi(x)$  may be huge/infinite).  
**Instead**, you use only dot products to make predictions!
  - Equ(decision function):

$$f(x) = \sum_{i=1}^m \hat{\alpha}(i) t(i) K(x(i), x) + \hat{b}$$

Only involves **kernel evaluations**, no explicit transformation.

- Final Note:
  - For very large-scale nonlinear problems, **Random Forests** or **Neural Networks** can sometimes be better choices than SVMs.

## Resources :

•

## Related notes :

•

---

# References :

- **Internal :**

- 
- 
- 

- **External :**

- [hegab videos](#)
- [the book](#)
- [the notebook](#)
- 

---

## 1. **Without scaling** (left plot in Figure 5-2):

- Suppose one feature (like  $x_0$ ) is between **0 and 6**, but another feature ( $x_1$ ) is between **0 and 80**.
- The SVM looks for the widest street *based on these raw values*.
- Because  $x_1$  is so much bigger than  $x_0$ , the SVM **thinks  $x_1$  is more important**, and the decision boundary becomes almost horizontal — **biased**.

---

## **With scaling** (right plot in Figure 5-2):

- After scaling, both features have similar ranges (like -2 to +2).
- Now, the SVM treats  $x_0$  and  $x_1$  **equally**.
- As a result, the **street is more natural**, better separating the points.

---

### **In short:**

⚡ **Scaling fixes the balance between features** so SVM can find the *true* best boundary without being tricked by big numbers. ↩

## 2. Out-of-core learning means **training a model without loading all the data into RAM at once**.

● Instead, it **loads small chunks of the data little by little**, processes them, and updates the model gradually.

→ This way, you can **train on very large datasets** that **don't fit into memory**.

---

**Simple example:**

- Your data is 500 GB.
- Your RAM is only 8 GB.
- → Out-of-core learning reads, say, 100 MB at a time, trains on it, then moves to the next 100 MB.

←