

2025-07-14 23:16

tags :

- [Machine Learning](#)
- [Hands on ML - book](#)

Chapter 6 - Hands on Training and Visualizing a Decision Tree

- A **DecisionTreeClassifier** is trained on the Iris dataset using two features: *petal length* and *petal width*.
- Example :

```
tree_clf = DecisionTreeClassifier(max_depth=2, random_state=42)
tree_clf.fit(X_iris, y_iris)`
```

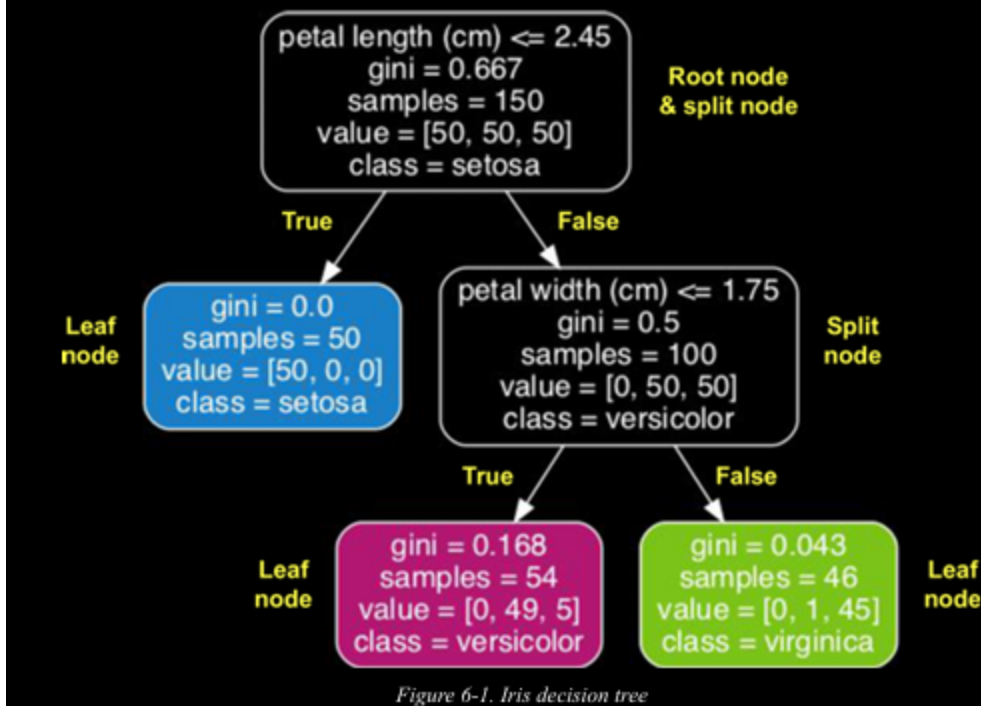
- Visualization:
 - Use `export_graphviz()` to create a `.dot` file for the tree.
 - Display in a Jupyter notebook using `graphviz.Source.from_file()`.

```
export_graphviz( tree_clf, out_file="iris_tree.dot",
                 feature_names=["petal length (cm)", "petal width (cm)"],
                 class_names=iris.target_names, rounded=True, filled=True
                 )
```

```
from graphviz import Source
Source.from_file("iris_tree.dot")
```

- Graphviz is an open source graph visualization software package. It also includes a dot command-line tool to convert `.dot` files to a variety of formats, such as PDF or PNG.

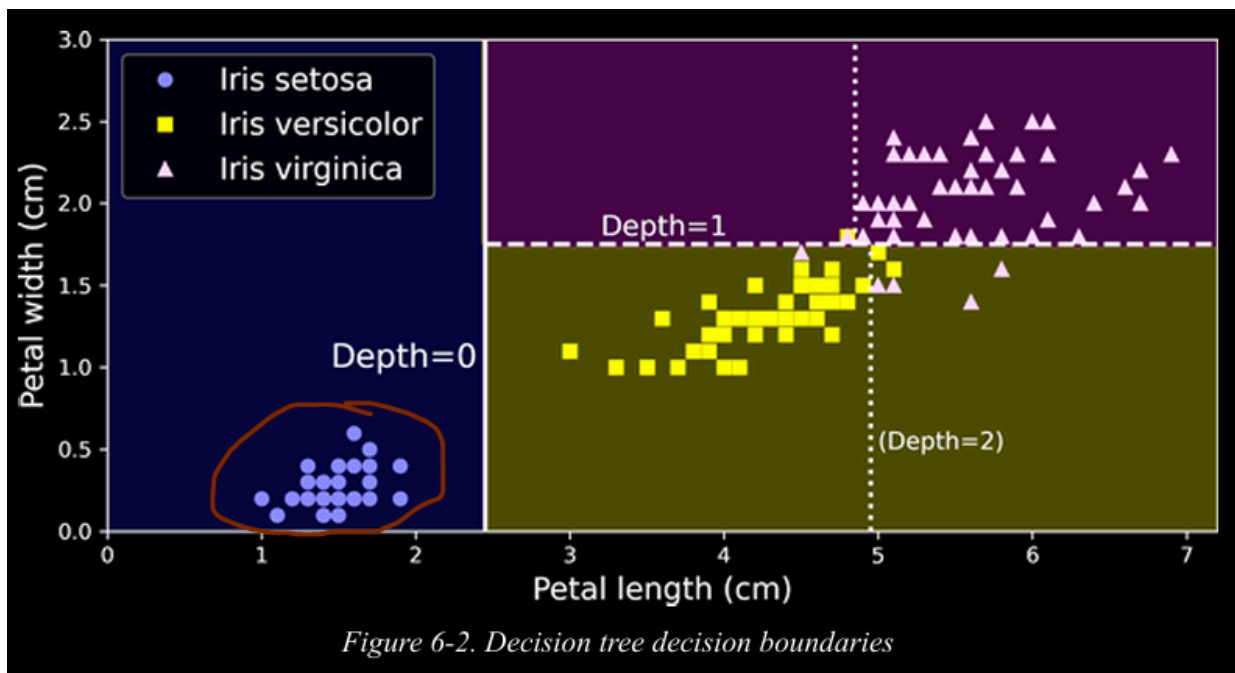
Your first decision tree looks like Figure 6-1.



Making Predictions

- To classify a flower, start at the root and follow the splits:
 - If **petal length ≤ 2.45 cm** → predict *Iris setosa* (leaf node).
 - Else → check **petal width ≤ 1.75 cm**:
 - If yes → predict *Iris versicolor*.
 - If no → predict *Iris virginica*.
- Advantages:**
 - No need for feature scaling or centering; trees work directly on raw data.- To classify a flower, start at the root and follow the splits:
 - If **petal length ≤ 2.45 cm** → predict *Iris setosa* (leaf node).
 - Else → check **petal width ≤ 1.75 cm**:
 - If yes → predict *Iris versicolor*.
 - If no → predict *Iris virginica*.
- Key Node Attributes:**
 - samples** : Number of training instances at the node.
 - value** : Count of instances per class at the node.
 - gini** : Measures impurity (0 = pure).
 - Example: Gini impurity at a node:

$$G = 1 - \sum_{k=1}^n (p_{i,k})^2$$
 where $p_{i,k}$ = ratio of class k in node i.
- Decision Tree Boundaries



- **CART algorithm** (used in Scikit-Learn) creates **binary trees**: each split produces **two branches** (yes/no decisions).
- At **depth 0**, the root splits data at *petal length* = 2.45 cm.
 - Left region is **pure** (*Iris setosa*) → no further splits.
 - Right region is **impure** → splits again at *petal width* = 1.75 cm (depth 1).
- With `max_depth=2`, the tree stops here. Increasing `max_depth` would add more splits and decision boundaries.
- **White Box vs Black Box Models**
- **Decision trees** are *white box models*: their decisions are **intuitive and interpretable**.
 - **Random forests & neural networks** are *black box models*: they often make accurate predictions but are **hard to interpret**.
 - **Interpretable ML** is an emerging field to help explain complex models' decisions (important for fairness and transparency).

Estimating Class Probabilities

- A decision tree can estimate the probability that an instance belongs to class k :
 1. Traverse the tree to find the **leaf node** for the instance.
 2. Compute the **ratio of class k instances** in that node.
- **Example**: For a flower with petals (5 cm, 1.5 cm) →
 - Probabilities:
 - *Iris setosa*: **0%** (0/54)
 - *Iris versicolor*: **90.7%** (49/54)
 - *Iris virginica*: **9.3%** (5/54)
 - Prediction: *Iris versicolor* (highest probability).

```
tree_clf.predict_proba([[5, 1.5]]) # Output: [[0.    , 0.907, 0.093]]
tree_clf.predict([[5, 1.5]]) # Output: [1]`
```

- ⚠ Limitation: Probabilities are **constant within each region**, even if some points seem closer to another class.

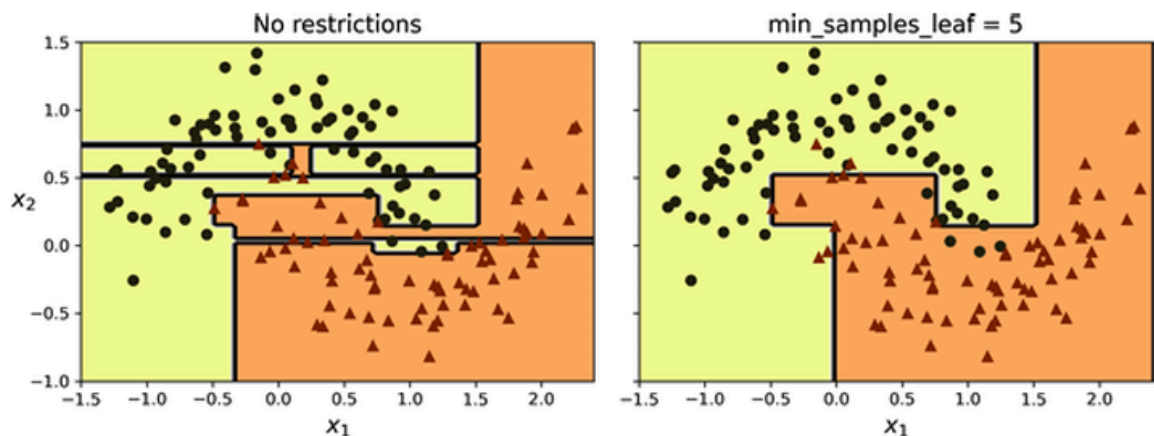
The CART Training Algorithm

- Scikit-Learn uses **CART (Classification and Regression Tree)** to build decision trees.
- At each step, it:
 1. Splits the data into two subsets using a feature `k` and threshold `t` (e.g., *petal length* ≤ 2.45 cm).
 2. Chooses (k, t) that minimizes the **cost function**:
$$J(k, t) = \frac{m_{\text{left}}}{m} G_{\text{left}} + \frac{m_{\text{right}}}{m} G_{\text{right}}$$
 - $G_{\text{left/right}}$: impurity of the subsets.
 - $m_{\text{left/right}}$: number of instances in each subset.
- **Algorithm Behavior**
 - **Recursively splits** the dataset based on best (k, t) pairs.
 - Stops when:
 - Maximum depth is reached (`max_depth`)
 - No split reduces impurity
 - Other stopping conditions triggered:
 - `min_samples_split`
 - `min_samples_leaf`
 - `min_weight_fraction_leaf`
 - `max_leaf_nodes`
- ⚠ **Limitations**
 - CART is a **greedy algorithm**: it looks for the best immediate split without considering future impact.
 - **Not optimal**: finding the globally best tree is **NP-complete** and computationally infeasible.
 - So CART aims for a "**reasonably good**" tree, not the best possible one.

Regularization Hyperparameters

- **Decision Trees** are nonparametric models:

- They don't have a fixed number of parameters.
- They adapt closely to the training data, which can lead to **overfitting**.
- To reduce overfitting, **regularization** limits the tree's flexibility.
- **Key Regularization Hyperparameters (Scikit-Learn):**
 - `max_depth` : Maximum depth of the tree (default = unlimited).
 - `max_features` : Max features to consider for splits at each node.
 - `max_leaf_nodes` : Limits number of leaf nodes.
 - `min_samples_split` : Minimum samples required to split a node.
 - `min_samples_leaf` : Minimum samples required to create a leaf node.
 - `min_weight_fraction_leaf` : Like `min_samples_leaf` but as a fraction of total weights.
- **Tip:** Increasing `min_*` or reducing `max_*` → more regularization.
- **Pruning** (used in some algorithms):
 - Start with a fully grown tree, then **remove unnecessary nodes**.
 - A node is pruned if its improvement is **not statistically significant** (tested with p-values, e.g., chi-squared test).
 - Pruning stops when all such nodes are deleted.
- **Testing Regularization on the Moons Dataset**
 - Two decision trees were trained:
 1. **Unregularized** (default settings).
 2. **Regularized** with `min_samples_leaf=5`.
- **Result (Figure 6-3):**



• *Figure 6-3. Decision boundaries of an unregularized tree (left) and a regularized tree (right)*

- The unregularized tree **overfits** the training data (complex decision boundary).
- The regularized tree produces a **simpler boundary** and generalizes better.
- **Test Set Accuracy:**
 - Unregularized tree: **89.8%**
 - Regularized tree: **92%** ✓

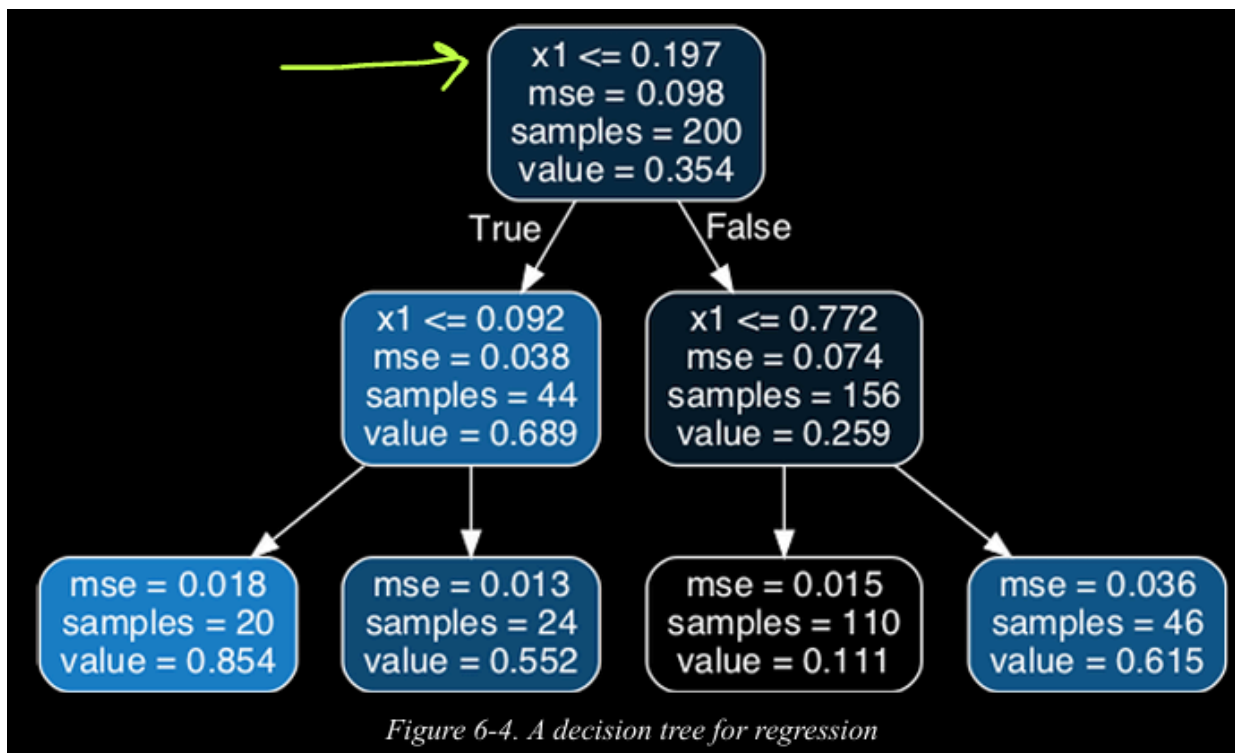
- **Key takeaway:** Regularization (e.g., `min_samples_leaf`) improves model generalization by reducing overfitting.

Regression

- Decision trees can also perform **regression tasks** using Scikit-Learn's `DecisionTreeRegressor`.
- Example: A **noisy quadratic dataset** is created, and a regression tree is trained with `max_depth=2`.

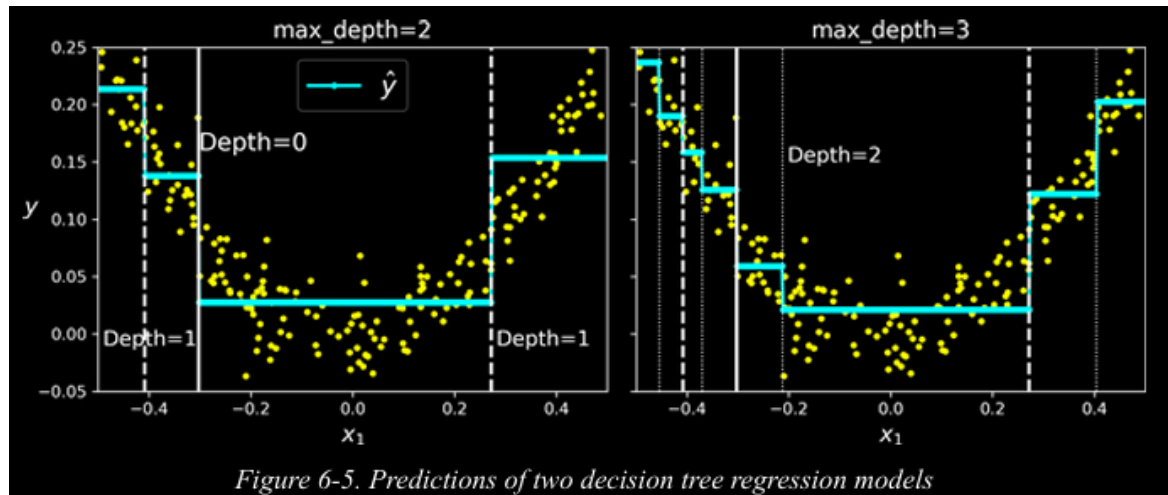
```
tree_reg = DecisionTreeRegressor(max_depth=2, random_state=42)
tree_reg.fit(X_quad, y_quad)`
```

- In the resulting tree (Figure 6-4), the splits are similar to those in classification, but instead of predicting classes, each **leaf node predicts a numerical value**.



- **How Prediction Works:**
 - For a new instance $x=0.2$:
 1. The root node asks if $x \leq 0.197$. Since $x = 0.2 > 0.197$, the algorithm moves to the **right child node**.
 2. Next, it asks if $x \leq 0.772$. Since $x = 0.2 < 0.772$, it moves to the **left child node**.
 3. This leaf predicts $\text{value}=0.111$, which is the **average target value** of the 110 training instances in that region.
 - The **Mean Squared Error (MSE)** of this prediction over those 110 instances is 0.015.
- **CART for Regression**

- The CART algorithm works similarly as in classification, but instead of minimizing impurity, it **minimizes MSE** at each split:



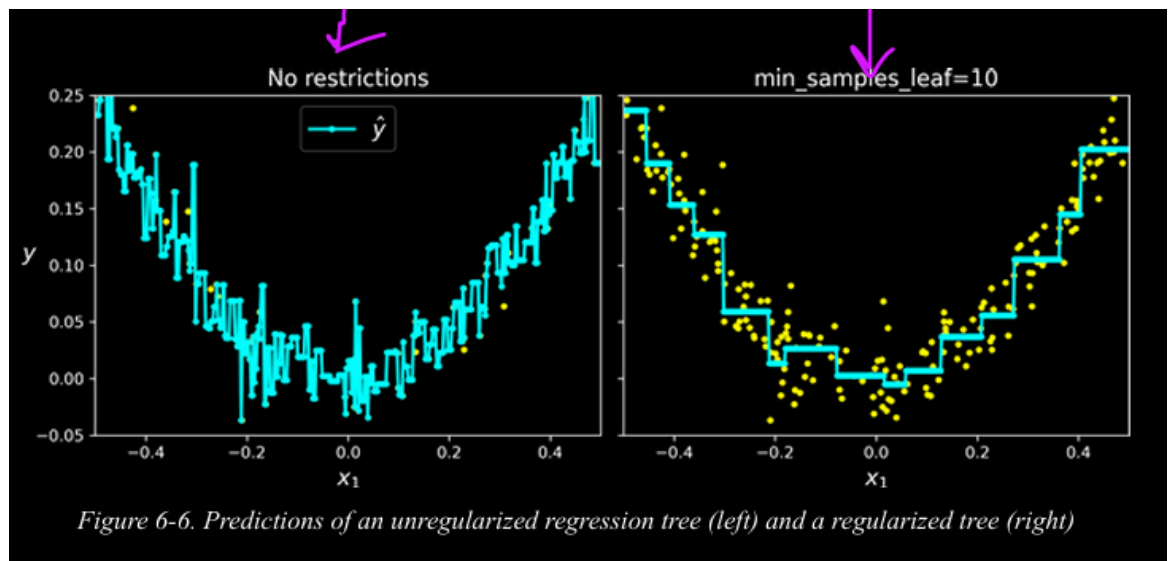
$$J(k, t_k) = \frac{m_{\text{left}}}{m} MSE_{\text{left}} + \frac{m_{\text{right}}}{m} MSE_{\text{right}}$$

where:

- $MSE_{\text{node}} = \frac{1}{m_{\text{node}}} \sum_{i \in \text{node}} (y^{(i)} - \hat{y}_{\text{node}})^2$
- \hat{y}_{node} : Mean target value in the node.

⚠ Overfitting in Regression Trees:

- Like in classification, regression trees are prone to **overfitting**, especially without regularization.
- Example (Figure 6-6):



- **Unregularized tree** (left): fits training data too closely, leading to poor generalization.
- **Regularized tree** (right): setting `min_samples_leaf=10` produces **smoother predictions** that generalize better.

Sensitivity to Axis Orientation

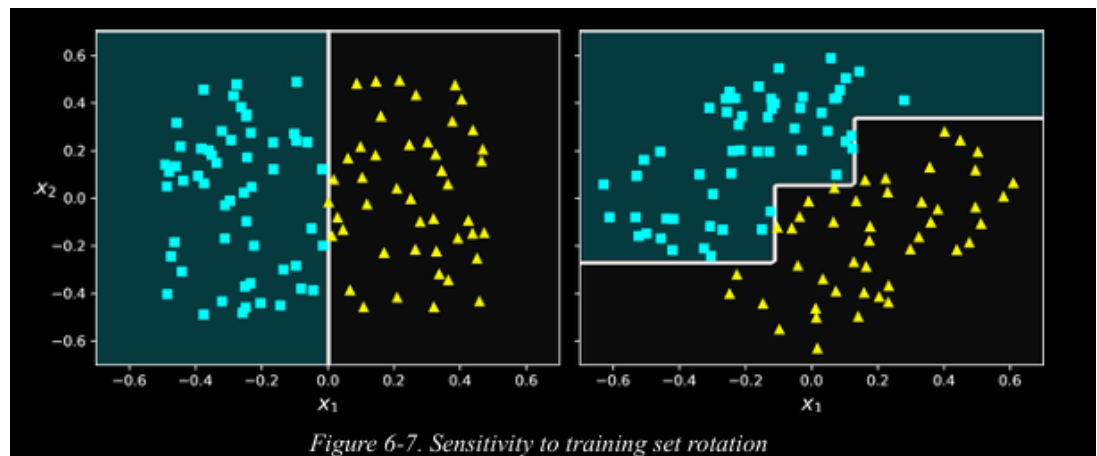
- **Strengths of Decision Trees:**

- Easy to understand and interpret.
- Simple to use, versatile, and powerful.

- **Limitation: Sensitivity to Orientation**

- Decision trees naturally create **orthogonal decision boundaries** (splits perpendicular to feature axes).
- This makes them **sensitive to the dataset's orientation**:

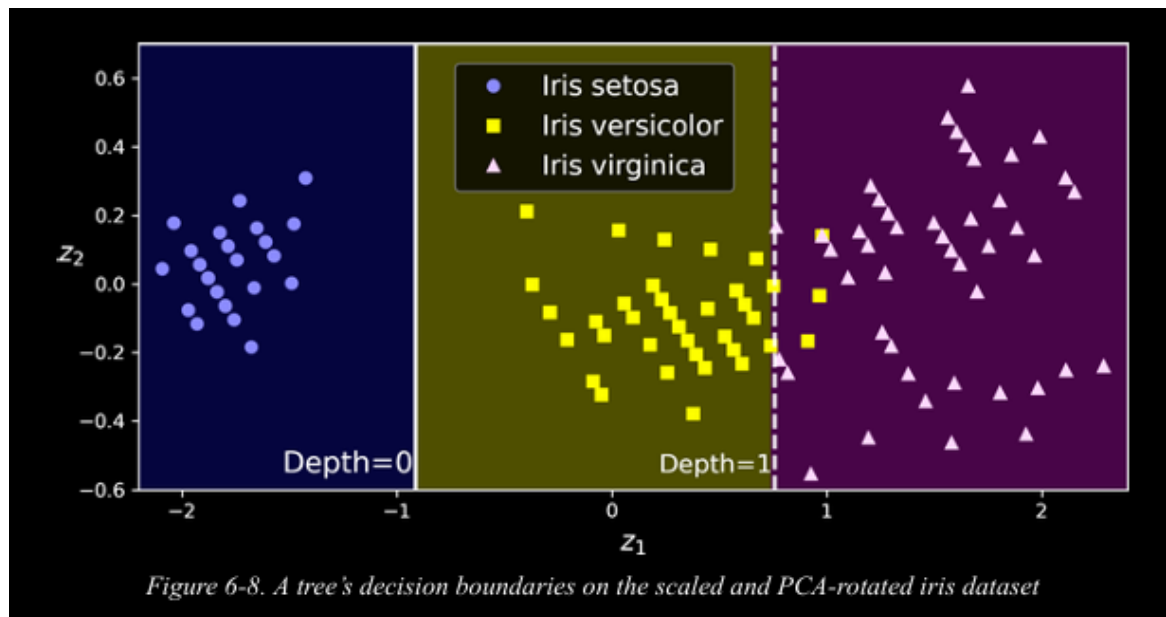
- 📊 *Example (Figure 6-7):*



- On the **left**, the dataset is aligned with the axes → the tree splits cleanly.
- On the **right**, the same dataset rotated by 45° results in a **complex and jagged decision boundary**, even though both trees perfectly fit the data.
- ⚠ The rotated tree is more likely to **overfit** and generalize poorly.

- **Solution: Scaling + PCA**

- To mitigate orientation sensitivity:
 1. **Scale the data** using `StandardScaler`.
 2. Apply **Principal Component Analysis (PCA)** to rotate the dataset and reduce feature correlations.
- *Example (Figure 6-8):*



- After scaling and PCA rotation, the Decision Tree fits the dataset well using only a single principal component (z_1), a linear combination of petal length and width.
- Code snippet for pipeline:

```
from sklearn.decomposition import PCA
from sklearn.pipeline import make_pipeline
from sklearn.preprocessing import StandardScaler

pca_pipeline = make_pipeline(StandardScaler(), PCA())
X_iris_rotated = pca_pipeline.fit_transform(X_iris)

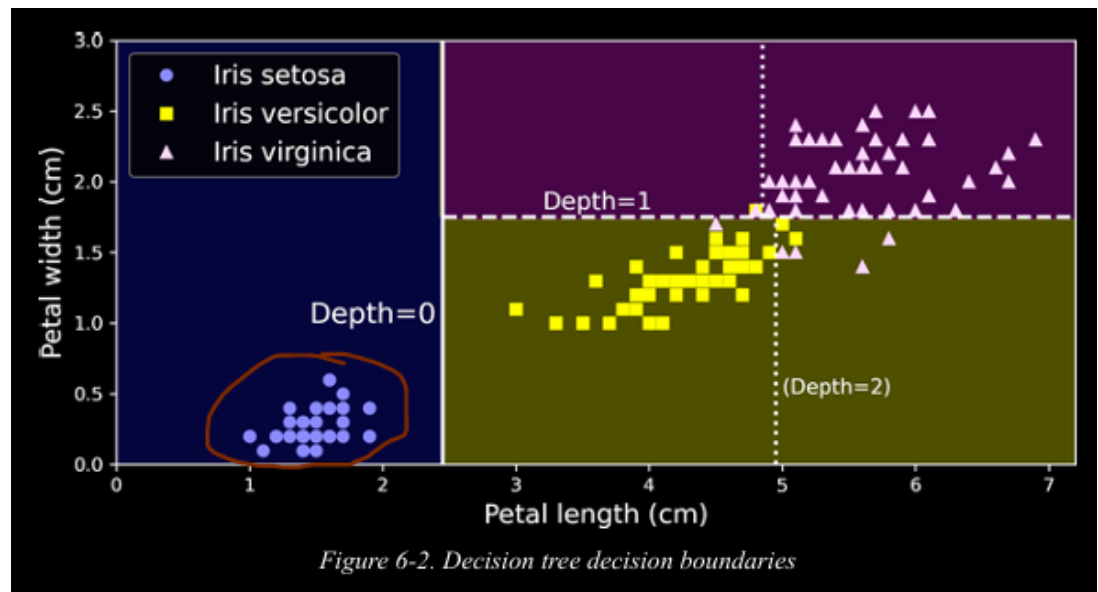
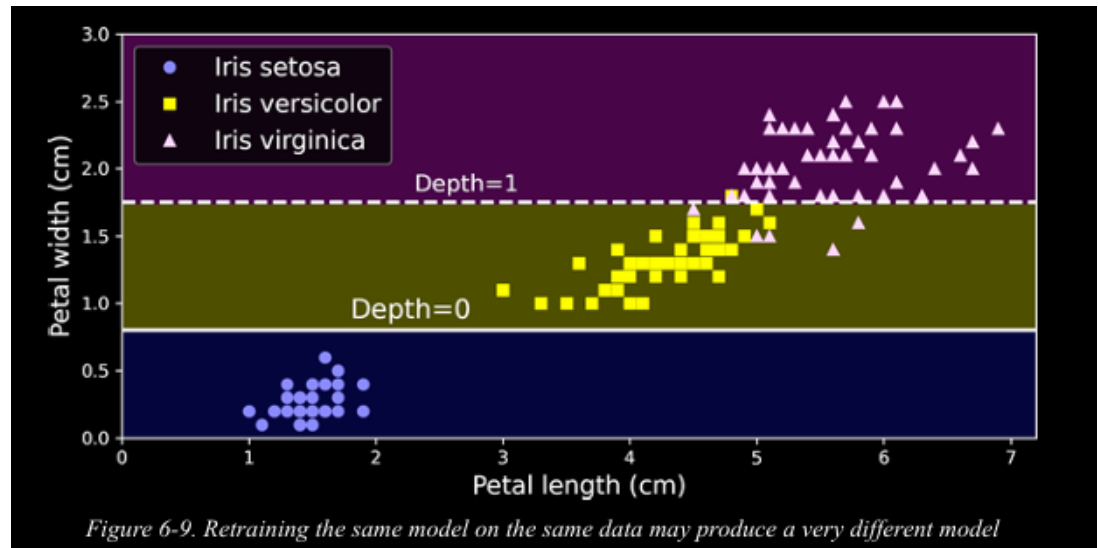
tree_clf_pca = DecisionTreeClassifier(max_depth=2, random_state=42)
tree_clf_pca.fit(X_iris_rotated, y_iris)
```

- **Key Takeaway:**
 - While decision trees are powerful, their reliance on axis-aligned splits can lead to poor performance on rotated datasets.
 - Preprocessing with **scaling and PCA** can improve their ability to create simpler, more generalizable decision boundaries.

Decision Trees Have a High Variance

- **Decision Trees Have High Variance**
 - A key limitation of decision trees is their **high variance**:
 - Small changes in the **training data** or **hyperparameters** can result in very different tree structures.

- Even retraining on the **same dataset** may produce a different tree because Scikit-Learn's training algorithm is **stochastic** (it randomly selects which features to evaluate at each node).
- Example: *Figure 6-9* shows a retrained tree that looks very different from the earlier tree in *Figure 6-2*.



- This variability can lead to **instability** in predictions.

• **Solution: Reduce Variance with Ensembles**

- A powerful way to address high variance is to **combine multiple trees** into an ensemble.
- Averaging predictions from many trees reduces variance significantly.
- ☒ Such an ensemble is called a **Random Forest**, which combines the predictions of multiple trees for better stability and performance.

 *Random Forests* will be covered in detail in **Chapter 7**.

Resources :

-
-

Related notes :

-
-

References :

- **Internal :**

-
-
-

- **External :**

- [hegab videos](#)
- [the book](#)
- [the notebook](#)
-