

2025-08-16 00:51

tags :

- [Machine Learning](#)
- [Hands on ML - book](#)

Chapter 7 - Hands on

Voting Classifiers

- You train several different classifiers (e.g., logistic regression, SVM, random forest, k-nearest neighbors).
- Each model has similar performance (around 80% accuracy).
- Instead of relying on a single model, you **combine their predictions**.
- The final prediction is the **class chosen by majority vote** among the models.
- This technique is called a **hard voting classifier**, and it often performs better than the individual models alone.
- A **voting classifier** can often outperform the best individual classifier in the ensemble.
- Even weak learners (slightly better than random guessing) can combine into a strong learner if there are enough of them and they are diverse.
- Analogy: tossing a biased coin (51% heads). With many tosses, the probability of a majority of heads increases (law of large numbers). Similarly, many weak classifiers can yield high accuracy.
- **Limitation:** if classifiers are correlated (make similar mistakes), the benefit is reduced.
- **Key tip:** Ensemble methods work best when classifiers are **independent and diverse** (trained with different algorithms).

voting Classifier in Scikit-Learn

- **Class:** `VotingClassifier` (easy to use → give list of (name, estimator) pairs).
- **Behavior:** Clones and fits all estimators.
 - Original models → `.estimators`
 - Fitted clones → `.estimators_`
 - Dict access → `.named_estimators_`

```
from sklearn.datasets import make_moons
from sklearn.ensemble import RandomForestClassifier, VotingClassifier
```

```

from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
from sklearn.svm import SVC
X, y = make_moons(n_samples=500, noise=0.30, random_state=42)
X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=42)
voting_clf = VotingClassifier( estimators=[
    ('lr', LogisticRegression(random_state=42)),
    ('rf', RandomForestClassifier(random_state=42)),
    ('svc', SVC(random_state=42)) ] )
voting_clf.fit(X_train, y_train)

```

Hard Voting

- Predicts class by **majority vote**.
- Example (moons dataset):
 - Logistic Regression = 86.4%
 - Random Forest = 89.6%
 - SVM = 89.6%
 - Voting Classifier = **91.2%** (better than all individuals).

Soft Voting

- Uses **average of predicted probabilities** → class with highest average prob.
- Usually performs better than hard voting.
- Requirement: all classifiers must support `predict_proba()`.
 - For SVM(not have `predict_proba()`), set `probability=True`.
 - this will make the SVC class use cross validation to estimate class probabilities, slowing down training, and it will add a `predict_proba()` method)
- Example: Soft voting improved accuracy to **92%**.

Key takeaway:

- Voting ensembles often outperform individual models, with **soft voting** usually being the best option if probability estimates are available.

Bagging and Pasting

- **Idea:** Use the same algorithm but train on **different random subsets** of the training set.
 - **Bagging (Bootstrap Aggregating):** sampling **with replacement**.
 - **Pasting:** sampling **without replacement**.
- **Sampling:**

- Across predictors → instances can repeat.
- Within a predictor → only Bagging allows repetition.
- **Prediction Aggregation**
 - **Classification:** mode (majority vote).
 - **Regression:** average.
- **Effect on Bias/Variance**
 - Each predictor individually → higher bias.
 - Aggregated ensemble →
 - **Bias \approx single predictor(fitted on the all dataset)**
 - **Variance \downarrow (reduced)** → better generalization.
- **Parallelism**
 - Training and prediction can be done **in parallel** (different cores/servers).
 - Makes Bagging & Pasting **scalable and efficient**.
- **Key takeaway:**
 - Bagging (with replacement) and Pasting (without replacement) create ensembles that reduce variance, often outperforming a single model, and are highly parallelizable.

Bagging and Pasting in Scikit-Learn

- **Classes:**
 - `BaggingClassifier` (classification)
 - `BaggingRegressor` (regression)
- **Key Parameters**
 - `n_estimators` : number of base models (e.g., 500 trees).
 - `max_samples` : number of training instances per base model (e.g., 100).
 - `bootstrap=True` → **Bagging** (with replacement).
 - `bootstrap=False` → **Pasting** (without replacement).
 - `n_jobs=-1` → use all CPU cores (parallel training/prediction).
- Code

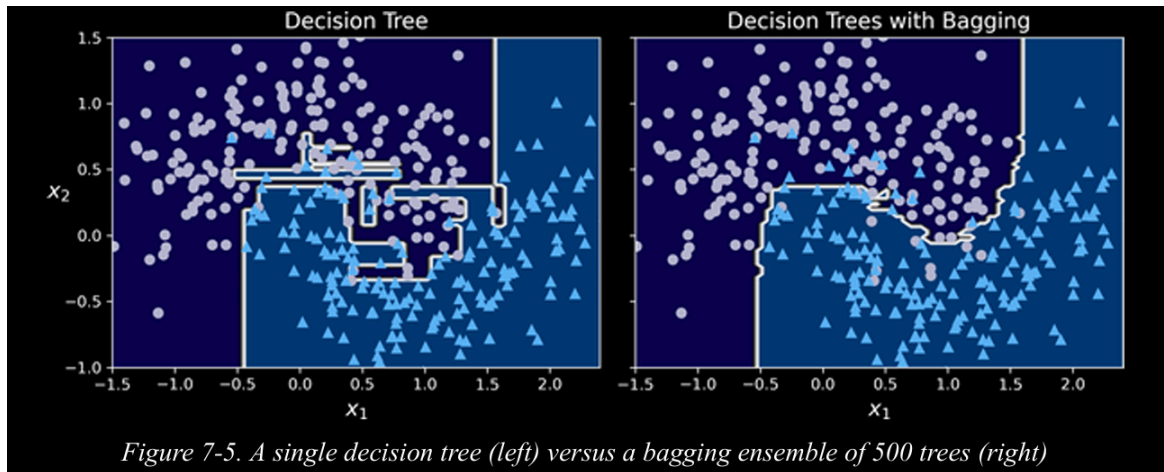
```
from sklearn.ensemble import BaggingClassifier
from sklearn.tree import DecisionTreeClassifier
bag_clf = BaggingClassifier(DecisionTreeClassifier(),
    n_estimators=500, max_samples=100, n_jobs=-1, random_state=42)
bag_clf.fit(X_train, y_train)
```

- **Soft Voting**
 - If base estimator supports `predict_proba()` (e.g., decision trees), ensemble uses

soft voting automatically.

- **Bias–Variance Effect**

- Bagging & Pasting → reduce **variance** compared to a single model.
- Bagging → slightly higher bias than pasting (more randomness), but **lower correlation** → usually better performance.
- **Generalization**: Ensemble smoother decision boundaries, less overfitting than a single decision tree.



- **Practical Tip**

- Bagging usually preferred.
- But test both (via cross-validation) if resources allow.

- **Key takeaway:**

`BaggingClassifier` (or `Regressor`) = train multiple predictors on random subsets of data. Bagging (with replacement) often beats pasting due to reduced correlation, giving better generalization.

Out-of-Bag Evaluation

- In **Bagging**:

- Each predictor is trained on a bootstrap sample (size = training set).
- On average → ~63% of training instances are sampled.
- Remaining ~37% = **Out-of-Bag (OOB) instances** (different for each predictor).

- **Use of OOB Instances**

- Serve as a **built-in validation set** → no need for a separate one.
- Each training instance is OOB for some predictors → use those to get ensemble predictions.
- Compute accuracy (or other metrics) based on OOB predictions.

- In **Scikit-Learn**

- Enable with `oob_score=True`.
- Accuracy available via `.oob_score_`.

- Class probability estimates available via `.oob_decision_function_`.

```
>>> bag_clf = BaggingClassifier(DecisionTreeClassifier(),
                               n_estimators=500, ... oob_score=True,
                               n_jobs=-1, random_state=42)
...
>>> bag_clf.fit(X_train, y_train)
>>> bag_clf.oob_score_ # 0.896

>>> from sklearn.metrics import accuracy_score
>>> y_pred = bag_clf.predict(X_test)
>>> accuracy_score(y_test, y_pred) # 0.92
```

- **Example Results**

- OOB score \approx **89.6%**.
- Actual test set accuracy = **92%** (OOB slightly pessimistic, ~2% lower).
- `.oob_decision_function_` → gives predicted class probabilities per instance.

- **Key takeaway:**

OOB evaluation provides a **free, internal validation estimate** for bagging models—usually close to test accuracy, but slightly pessimistic.

Random Patches and Random Subspaces

- **Feature Sampling in Bagging**

- Controlled by:
 - `max_features` → number/ratio of features to sample.
 - `bootstrap_features` → whether sampling is **with replacement**.
- Works like `max_samples` & `bootstrap`, but for **features** instead of instances.

- **Methods**

1. **Random Patches**

- Sample both **instances** and **features**.
- Default bagging + feature sampling.

2. **Random Subspaces**

- Keep **all training instances** (`bootstrap=False`, `max_samples=1.0`).
- Sample only **features** (`bootstrap_features=True` and/or `max_features < 1.0`).

- **Effect**

- Increases **diversity** among predictors.
- Leads to **lower variance**, but adds a bit more **bias**.
- Particularly useful for **high-dimensional data** (e.g., images) → speeds up training.

- **Key takeaway:**

Feature sampling (random patches/subspaces) enhances diversity and reduces variance, at the cost of slightly higher bias—especially helpful in high-dimensional problems.

Random Forests

- **Random Forests**

- **Definition:** Ensemble of decision trees, usually trained with **bagging** (sometimes pasting).

- **Classes:**

- `RandomForestClassifier` (classification)
- `RandomForestRegressor` (regression)

- **Convenience:** Has all hyperparameters of

- `DecisionTreeClassifier` (tree growth control)
- `BaggingClassifier` (ensemble control).

- **Extra Randomness**

- At each split, instead of checking all features → use a **random subset**.
- Default = \sqrt{n} features (classification).
- Increases **diversity** among trees.
- Effect:
 - Slightly ↑ bias
 - ↓ variance
 - Better overall generalization.

- **Equivalence**

```
RandomForestClassifier(n_estimators=500, max_leaf_nodes=16)
```

≈




```
BaggingClassifier(      DecisionTreeClassifier(max_features="sqrt",  
max_leaf_nodes=16),      n_estimators=500 )
```

- **Key takeaway:**

Random Forests = bagging of decision trees + **feature randomness** at splits → strong, low-variance models with good generalization.

Extra-Trees

- Extra-Trees (Extremely Randomized Trees)

- In **Random Forests**, each split considers a random subset of features.
- In **Extra-Trees**, extra randomness is added:
 - **Random thresholds** are chosen for splits (instead of searching for the best threshold).
- **Trade-off:**
 -  Lower variance
 -  Faster training (no need to search best thresholds)
 -  Higher bias
- **Implementation in Scikit-Learn:**
 - `ExtraTreesClassifier` (for classification)
 - `ExtraTreesRegressor` (for regression)
 - API is the same as `RandomForest` classes, but **bootstrap=False by default**.
- **Key Point:** No way to know in advance if Random Forests or Extra-Trees work better → must test both with **cross-validation**.

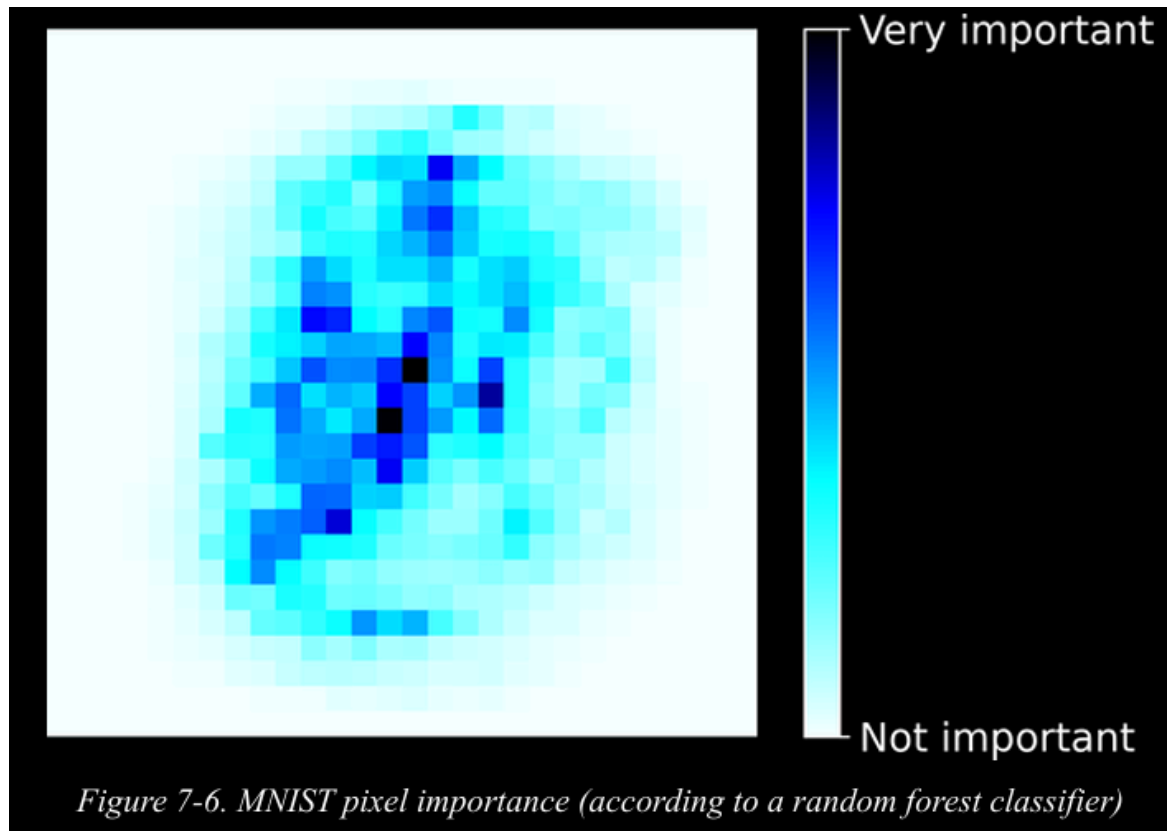
Feature Importance

- Feature Importance in Random Forests
 - **Definition:** Importance of a feature is based on how much it reduces impurity across all trees.
 - **Weighting:** Each node's contribution is weighted by the number of training samples it affects.
 - **Scaling:** Importances are normalized so they sum to 1.
 - **Access:** Available through `feature_importances_` after training.
- **Example (Iris dataset):**

```
from sklearn.datasets import load_iris
iris = load_iris(as_frame=True)
rnd_clf = RandomForestClassifier(n_estimators=500, random_state=42)
rnd_clf.fit(iris.data, iris.target)
for score, name in zip(rnd_clf.feature_importances_, iris.data.columns):
    print(round(score, 2), name)
```

- Sepal length → **0.11**
 - Sepal width → **0.02**
 - Petal length → **0.44**
 - Petal width → **0.42**
- Key Uses:

- Helps **interpret models** (which features matter most).
- Useful for **feature selection**.
- Can visualize importance (e.g., MNIST pixel map).



Boosting

- **Definition:** Ensemble method that combines several **weak learners** into a **strong learner**.
- **Key Idea:** Learners are trained **sequentially**, each new model focuses on correcting the errors of the previous one.
- **Popular Methods:**
 - **AdaBoost** (Adaptive Boosting)
 - **Gradient Boosting**
- Purpose: Improve accuracy by reducing bias and variance compared to single weak learners.

AdaBoost

- Core idea
 - Build an ensemble **sequentially**: each new weak learner focuses more on **instances misclassified** by previous learners (by **increasing their weights**).
 - Final prediction is a **weighted vote** of all learners.
- Training procedure (binary or multiclass via SAMME)

- **Initialization**

- For mmm training instances, set **instance weights**:

$$w_i^{(1)} = \frac{1}{m} \text{ for } i = 1, \dots, m.$$

- **For $j = 1 \dots N$ (number of estimators):**

1. **Fit** base learner h_j on the training set **using current weights** $\{w_i^{(j)}\}$.

2. **Weighted error**:

$$r_j = \sum_{i=1}^m w_i^{(j)} \cdot \mathbf{1}[h_j(x_i) \neq y_i]$$

3. **Learner weight (vote)**:

$$\alpha_j = \eta \cdot \log\left(\frac{1-r_j}{r_j}\right)$$

- η = **learning rate** (default 1.0).
- If $r_j < 0.5 \rightarrow \alpha_j > 0$ (useful learner).
- If $r_j = 0.5 \rightarrow \alpha_j = 0$ (no influence).
- If $r_j > 0.5 \rightarrow \alpha_j < 0$ (worse than random).

4. **Update instance weights :**

$$w_i^{(j+1)} = \begin{cases} w_i^{(j)} & \text{if } h_j(x_i) = y_i \\ w_i^{(j)} \cdot \exp(\alpha_j) & \text{if } h_j(x_i) \neq y_i \end{cases}$$

- Then **normalize**: $w_i^{(j+1)} \leftarrow \frac{w_i^{(j+1)}}{\sum_{k=1}^m w_k^{(j+1)}}$

5. **Stop** when N learners are trained or a **perfect** learner is reached.

Intuition: Misclassified points **gain weight** by a factor $\exp(\alpha_j)$, pulling the next learner to focus on the “hard” cases. Smaller η (learning rate) \rightarrow smaller weight boosts each round (smoother learning, often better generalization).

- Prediction rule

- **Weighted majority vote (SAMME)**:

$$\hat{y}(x) = \arg \max_k \sum_{j=1}^N \alpha_j \cdot \mathbf{1}[h_j(x) = k]$$

- **SAMME.R** (if base learners output probabilities): uses **class probabilities** instead of hard labels; generally performs **better**.

- Relation to gradient descent

- Conceptually similar: instead of tweaking one model’s parameters to minimize loss, AdaBoost **adds learners** sequentially, each step reducing the ensemble’s error (see the “decision boundaries” plots: halving the learning rate \rightarrow gentler updates, smoother boundaries).

- Parallelism & scaling

- **Drawback:** Training is **sequential** (each round depends on the previous), so it **cannot be parallelized** like bagging/pasting.

- Scikit-Learn essentials

- **Classes:** `AdaBoostClassifier`, `AdaBoostRegressor`.

- **Default base estimator: decision stump**

`DecisionTreeClassifier(max_depth=1)` .

- **Common params:**

- `n_estimators` → number of weak learners NNN
- `learning_rate` → η (shrinks α_j)

- **Code pattern:**

```
from sklearn.ensemble import AdaBoostClassifier
from sklearn.tree import DecisionTreeClassifier
ada_clf = AdaBoostClassifier(
    DecisionTreeClassifier(max_depth=1),
    n_estimators=30,
    learning_rate=0.5,
    random_state=42 )
ada_clf.fit(X_train, y_train)
```

- **Multiclass:** uses **SAMME**; if base learners support `predict_proba` , Scikit-Learn can use **SAMME.R** (usually better).
- Practical tips
 - If **overfitting**:
 - Decrease `n_estimators` , and/or
 - **Regularize** the base learner more (e.g., shallower trees, min samples split/leaf).
 - Use smaller `learning_rate` with larger `n_estimators` for a stronger but smoother ensemble.
- Ultra-short flashcard
 - **Updates:** $\alpha_j = \eta \log \frac{1-r_j}{r_j}$, miscls weights $\times e^{\alpha_j}$, then **normalize**.
 - **Predict:** weighted vote by α_j .
 - **Pros:** strong from weak learners; **Cons:** sequential (no parallel).
 - **SAMME / SAMME.R** for multiclass; **stumps** are common base learners.

Gradient Boosting(GBRT)

- Core idea
 - Like AdaBoost, builds ensemble **sequentially**.
 - Instead of reweighting misclassified samples, each new predictor is trained to **fit the residual errors** of the previous predictors.
 - Base learners: usually **decision trees** → Gradient Boosted Regression Trees (GBRT).
- Step-by-step (regression example)
 1. Fit first tree $h_1(x)$ to data.

2. Compute residuals $r_1 = y - h_1(x)$.
3. Fit second tree $h_2(x)$ to residuals.
4. Compute new residuals $r_2 = r_1 - h_2(x)$.
5. Continue for N trees.

- **Final prediction:**

$$\hat{y}(x) = \sum_{j=1}^N h_j(x)$$

- Scikit-Learn implementation

```
from sklearn.ensemble import GradientBoostingRegressor
gbrt = GradientBoostingRegressor(
    max_depth=2,
    n_estimators=3,
    learning_rate=1.0,
    random_state=42 )
gbrt.fit(X, y)
```

- GradientBoostingRegressor (for regression)
- GradientBoostingClassifier (for classification)
- Key hyperparameters
 - **Tree growth controls:** max_depth, min_samples_leaf, etc.
 - **Ensemble controls:**
 - n_estimators : number of trees
 - learning_rate : scales contribution of each tree (smaller → need more trees, but better generalization → called **shrinkage**)
 - subsample : fraction of training instances per tree → stochastic gradient boosting (adds randomness, reduces variance, speeds training)
- Regularization & Overfitting control
 - **Shrinkage (learning_rate):**
 - Lower learning_rate + higher n_estimators → better generalization.
 - **Early stopping:**
 - n_iter_no_change : stop if no improvement for k iterations.
 - Uses an internal validation set (validation_fraction , default 10%).
 - tol : threshold for negligible improvement.
 - **Subsampling:**
 - Train each tree on a subset of training set (e.g., subsample=0.25) → higher bias, lower variance, faster training.
- Practical tips
 - Use **grid/randomized search** to find n_estimators and learning_rate .

- Safer default: small `learning_rate` (0.05–0.1) with large `n_estimators`.
- Early stopping (`n_iter_no_change`) helps pick optimal number of trees automatically.
- For very large datasets: consider stochastic gradient boosting (`subsample < 1.0`).
- Flashcard summary
 - **Trains on residuals**, not weights (unlike AdaBoost).
 - Prediction = sum of all trees.
 - **Shrinkage** = small `learning_rate` + more trees.
 - **Regularization**: early stopping, subsampling.
 - **Tradeoff**: more trees → better fit, but risk of overfitting.

Histogram-Based Gradient Boosting

- **What it is:**

An optimized implementation of Gradient Boosted Regression Trees (GBRT) in Scikit-Learn, designed for **large datasets**.
- **How it works:**
 - Bins input features into integers (`max_bins` hyperparameter, default = 255, max = 255).
 - Binning reduces the number of thresholds to evaluate.
 - Using integers → faster & memory-efficient data structures.
 - Removes need to **sort features** when training each tree.
- **Complexity:**
 - Standard GBRT: $O(n \times m \times \log(m))$
 - HGB: $O(b \times m)$
 - `n` : number of features
 - `m` : number of instances
 - `b` : number of bins
 - ⇒ HGB can be **hundreds of times faster** on large datasets.
- **Trade-off:**
 - Binning causes **precision loss**.
 - Acts as a **regularizer** → can reduce overfitting or may cause underfitting (dataset dependent).
- Classes in Scikit-Learn:
 - `HistGradientBoostingRegressor`
 - `HistGradientBoostingClassifier`

(similar to `GradientBoostingRegressor` / `Classifier` , with key differences).
- Key Differences from Standard GBRT:
 1. **Early Stopping**

- Auto-activated if dataset > 10,000 instances.
- Can force with `early_stopping=True/False`.

2. Subsampling

- **Not supported** in HGB.

3. `n_estimators` → renamed to `max_iter`

4. Tree hyperparameters allowed:

- `max_leaf_nodes`
- `min_samples_leaf`
- `max_depth`

5. Supports categorical features & missing values

- No need for imputer, scaler, or one-hot encoder.
- **Categorical features must be integers** (0 to < `max_bins`).
- Use `OrdinalEncoder` for transformation.
- Example Pipeline (California Housing Dataset)

```
from sklearn.pipeline import make_pipeline
from sklearn.compose import make_column_transformer
from sklearn.ensemble import HistGradientBoostingRegressor
from sklearn.preprocessing import OrdinalEncoder

hgb_reg = make_pipeline(
    make_column_transformer(
        (OrdinalEncoder(), ["ocean_proximity"]),
        remainder="passthrough"
    ),
    HistGradientBoostingRegressor(categorical_features=[0],
                                  random_state=42) )
hgb_reg.fit(housing, housing_labels)
```

- Very compact pipeline.
- No imputer, scaler, or one-hot encoder required.
- Without tuning → **RMSE ≈ 47,600**.
- Extra Note (TIP)
 - Other **optimized gradient boosting libraries**:
 - **XGBoost, CatBoost, LightGBM** → mature, GPU support, many features.
 - **TensorFlow Random Forests** → optimized implementations for random forests, Extra-Trees, GBRT, etc.

Stacking

- **Idea:** Instead of using simple aggregation (like voting), train a **model (meta-learner or blender)** to combine the predictions of base models.
- **Process:**
 1. Train base predictors.
 2. Use **cross_val_predict()** to generate out-of-sample predictions from each base model.
 3. These predictions form a **new training set** (each model gives 1 feature).
 4. Train the blender (final estimator) on this new dataset with the original targets.
 5. Retrain base models on the full dataset and use them + blender for final prediction.
- **Multi-layer stacking:**
 - Possible to train multiple blenders (different algorithms).
 - Outputs can feed into another blender (multi-layer stacking).
 - Can slightly improve performance but increases complexity and training cost.
- **Scikit-Learn Implementation**
 - Classes:
 - StackingClassifier
 - StackingRegressor
 - Example:

```
from sklearn.ensemble import StackingClassifier
stacking_clf = StackingClassifier(
    estimators=[
        ('lr', LogisticRegression(random_state=42)),
        ('rf', RandomForestClassifier(random_state=42)),
        ('svc', SVC(probability=True, random_state=42))
    ],

    final_estimator=RandomForestClassifier(random_state=43),
    cv=5)
stacking_clf.fit(X_train, y_train)
```

- Behavior:
 - Uses `predict_proba()` if available, else `decision_function()`, else `predict()`.
 - Default final estimator:
 - LogisticRegression for classifiers.
 - RidgeCV for regressors.
- **Performance & Takeaways**
- Stacking often improves accuracy (example: **92.8% vs 92% with soft voting**).
- **Best practices:**

- Try **Random Forests, AdaBoost, Gradient Boosted Trees** first (strong, versatile).
 - Ensembles are especially good for **heterogeneous tabular data**.
 - Require **little preprocessing**, making them great for fast prototyping.
 - Voting + Stacking can help **squeeze maximum performance**.
-

Resources :

-
-

Related notes :

-
-

References :

- **Internal :**

-
-
-

- **External :**

- [hegab videos](#)
- [the book](#)
- [the notebook](#)
-